

**docker**



# DOCKER I KUBERNETES – OD ZERA DO BOHATERA

Maciej Krajewski

Online – 18-22.11

KONFERENCJE ON-LINE SĄ JAK SEANSE  
SPIRYTYSTYCZNE

JOE MONSTER

GRAŻYNA JESTEŚ TAM?  
POWIEDZ COŚ, JEŚLI NAS SŁYSZYSZ  
JEST TAM KTOŚ Z TOBA?

# RULES



Write your name in zoom



If you are done with task – use zoom reactions



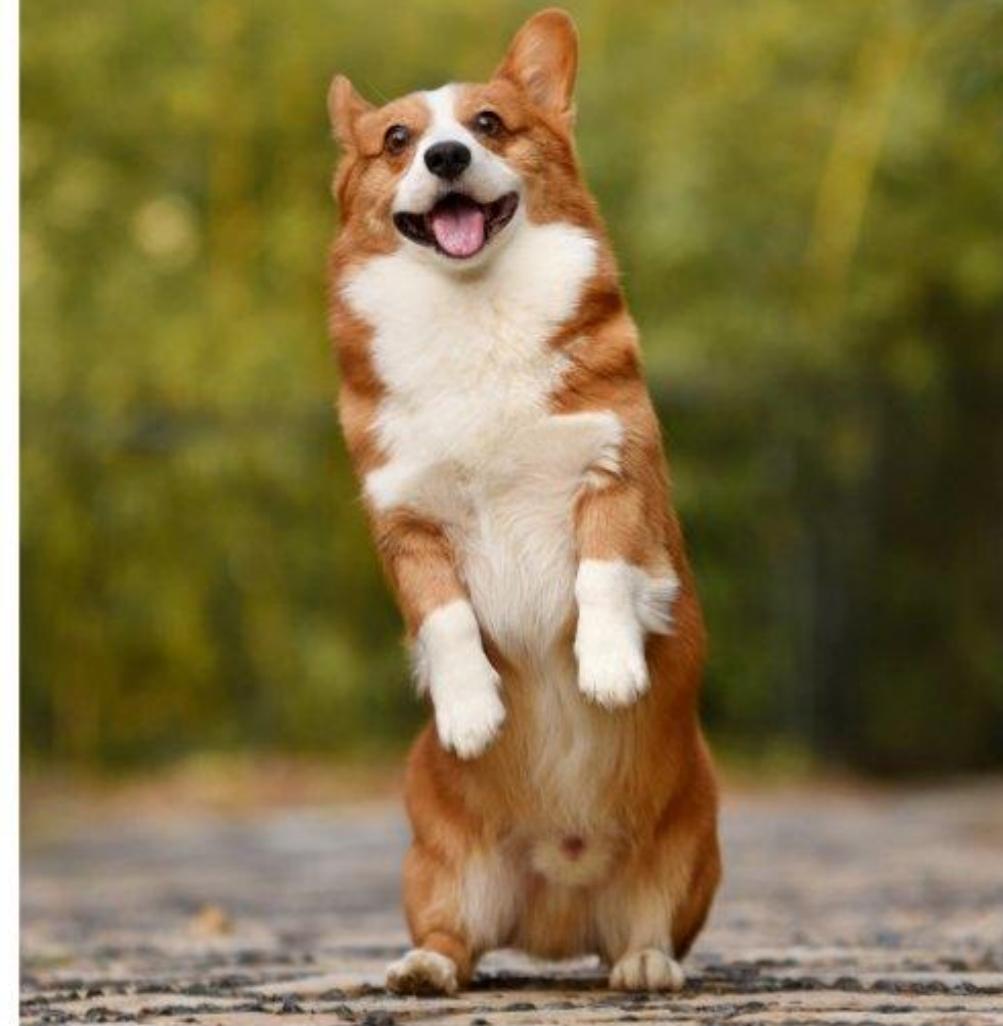
Stand up during breaks



Slack

**HI**

Hello there!



# BREAKS

## DAY 1-5

9:00 – 10:30

10:45 – 12:15

13:15 – 14:45

15:00 – 16:00

NOTE: OF COURSE ANY ORGANIZATION CAN USE THE LIGHTWEIGHT FORM (E.G. BY CREATING SELF-SUFFICIENT TEAMS)

# DOCUMENTATION

UX Knowledge Base Sketch #72



LIGHTWEIGHT

E.G.: AGILE + LEAN PROCESSES  
(BUT IN CASE OF THIRD PARTY DEVELOPMENT,  
YOU CAN'T DO EVERYTHING THE LEAN WAY)

IT CAN ALSO BE A MIX, A TRANSITION:  
AT THE BEGINNING OF THE PROJECT: LEAN,  
LATER: STANDARDIZED DOCUMENTATION

E.G.: ENTERPRISE-  
INTERNAL & REGULATORY  
COMPLIANCE



HEAVY WEIGHT



CLOSE COLLABORATION OF A CROSS-FUNCTIONAL TEAM  
EVERYONE IS ALREADY ON THE SAME PAGE.



"CUSTOM" CONTENT & STRUCTURE:  
• WHAT IS THE MOST EFFECTIVE WAY?  
• MINIMUM AMOUNT OF INFORMATION



IDEAL STAKEHOLDER MINDSET:  
FOCUSING ON THE OUTCOMES,  
IMPACT ACHIEVED BY THE PROJECT



SOME FORMS: SKETCHES (PAPER OR WHITEBOARD), PROTOTYPES (EVEN PAPER PROTOTYPES - YOU CAN EVOLVE IT BASED ON TEST RESULTS), DOCUMENTING COLLABORATIVE SESSIONS (E.G. CELLPHONE PHOTOS)



THE DOCUMENTATION CONSISTS OF LIVING DOCUMENTS - EDITED & UPDATED CONTINUALLY  
CO-CREATED DOCUMENTATION - FEEDBACK & QUESTIONS



IDEAL FOR ITERATIVE & INCREMENTAL PRODUCT DEVELOPMENT PROCESSES  
WORKING SOFTWARE OVER COMPREHENSIVE DOCUMENTATION<sup>4</sup>

AGILE MANIFESTO



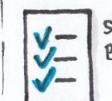
MAIN GOALS:  
• ELIMINATING WASTE  
• MOVING FORWARD FASTER  
• CAPTURE CONVERSATIONS (→ RECALL)



THE RELEASED PRODUCT/FEATURE ITSELF IS ALSO A DOCUMENTATION  
→ CONTINUOUS DELIVERY & IMPROVEMENT → HEAVY DOCUMENTATION WOULD QUICKLY BECOME OBSOLETE



IN MANY CASES:  
SILOS, ISOLATED TEAMS



STRICT STANDARDS & RULES  
BINDING CONTRACTS



USUAL STAKEHOLDER MINDSET:  
FOCUSING ON SHINY DESIGN ARTIFACTS & DELIVERABLES



SOME FORMS:  
REQUIREMENTS DOCUMENTS  
HUGE PILE OF DETAILED SPECIFICATION  
BIG DESIGN UP FRONT



THE DOCUMENTATION IS  
"CARVED IN STONE",  
IT'S HARDER TO MODIFY IT LATER IN THE PROCESS



APPROPRIATE FOR  
WATERFALL PRODUCT  
DEVELOPMENT PROCESSES WITH  
BIG HAND-OFFS



MAIN GOALS, REASONS:  
• THE PROCESS HEAVILY RELIES ON DOCUMENTATION  
• IT NEEDS TO BE UNDERSTANDABLE FOR OTHER TEAMS



DRAWBACKS:  
TAKES HUGE AMOUNT OF TIME & CREATES A LOT OF WASTE  
(BUT SOMETIMES IT IS NECESSARY)

## ADVICE



WHO IS THE TARGET AUDIENCE OF YOUR DOCUMENTATION? → DESIGN IT!

- FOR YOURSELF
- FOR YOUR TEAM
- FOR THE CLIENT
- FOR A THIRD PARTY TEAM

{ DIFFERENT GOALS  
MOTIVATION  
TERMINOLOGY }

E.G. FOR DEVELOPERS:  
• IMPLEMENTATION  
• DATA MODELING  
• DATA TYPES ETC.

{ UX OF THE DOCUMENTATION,  
E.G. FINDABILITY }

→ SINGLE SOURCE OF TRUTH  
DOCUMENTATION SHOULD BE EASILY ACCESSIBLE FOR THE WHOLE TEAM

• COLLABORATIVE: CONTRIBUTIONS, VERSIONING

• KEEP IT UP TO DATE

COMBINE WORDS & ILLUSTRATIONS (VISUALS)  
E.G. ANIMATED GIFS

CREATE A GRAVEYARD FOR DISCARDED VERSIONS, IDEAS



YOUR DESIGN PROCESS SHOULD DETERMINE WHAT THINGS YOU DOCUMENT, DOCUMENTATION SHOULD COMPLEMENT & ENHANCE THE PROCESS (E.G.: RESEARCH, SYNTHESIZING, IDEATION, DESIGN, TEST)



ORGANIZE LIVE DEMO SESSIONS,  
INVITE YOUR TEAM TO OBSERVE USABILITY TESTS  
→ FIRSTHAND INFORMATION IS THE BEST!

RESULT:  
NO DEPENDENCY  
ON THE DOCUMENTATION

INCLUDE WHAT PROBLEMS YOU ARE ADDRESSING, ADD CONTEXT, DISCUSSIONS, INSPIRATIONS...

## COMMON GOALS



• DOCUMENTING YOUR WAY OF THINKING  
• INFORMATION FOR YOUR FUTURE SELF / A TEAM / THIRD PARTIES



• CREATING A DECISION-MAKING TOOL  
(MORE INFORMED DECISIONS)



• TEAM ALIGNMENT,  
SHARED VISION  
• EXPLAINING THE "HOW" AND THE "WHY"

(AND SOMETIMES THE GOAL IS TO BE  
IN ACCORDANCE WITH A CONTRACT)

## + ANOTHER EXAMPLES:

AGENCIES & FREELANCERS WHO'RE GETTING PAID FOR CREATING DOCUMENTATIONS.  
IF THEY WANT TO SWITCH TO LIGHTWEIGHT, FREQUENT CLIENT INVOLVEMENT CAN HELP.

TRY FAIL SUCCESS



# AGENDA

Microservices

Rest API

DOCKER

KUBERNETES

Monitoring Aplikacji

Centralne Logowanie

12 Factor App

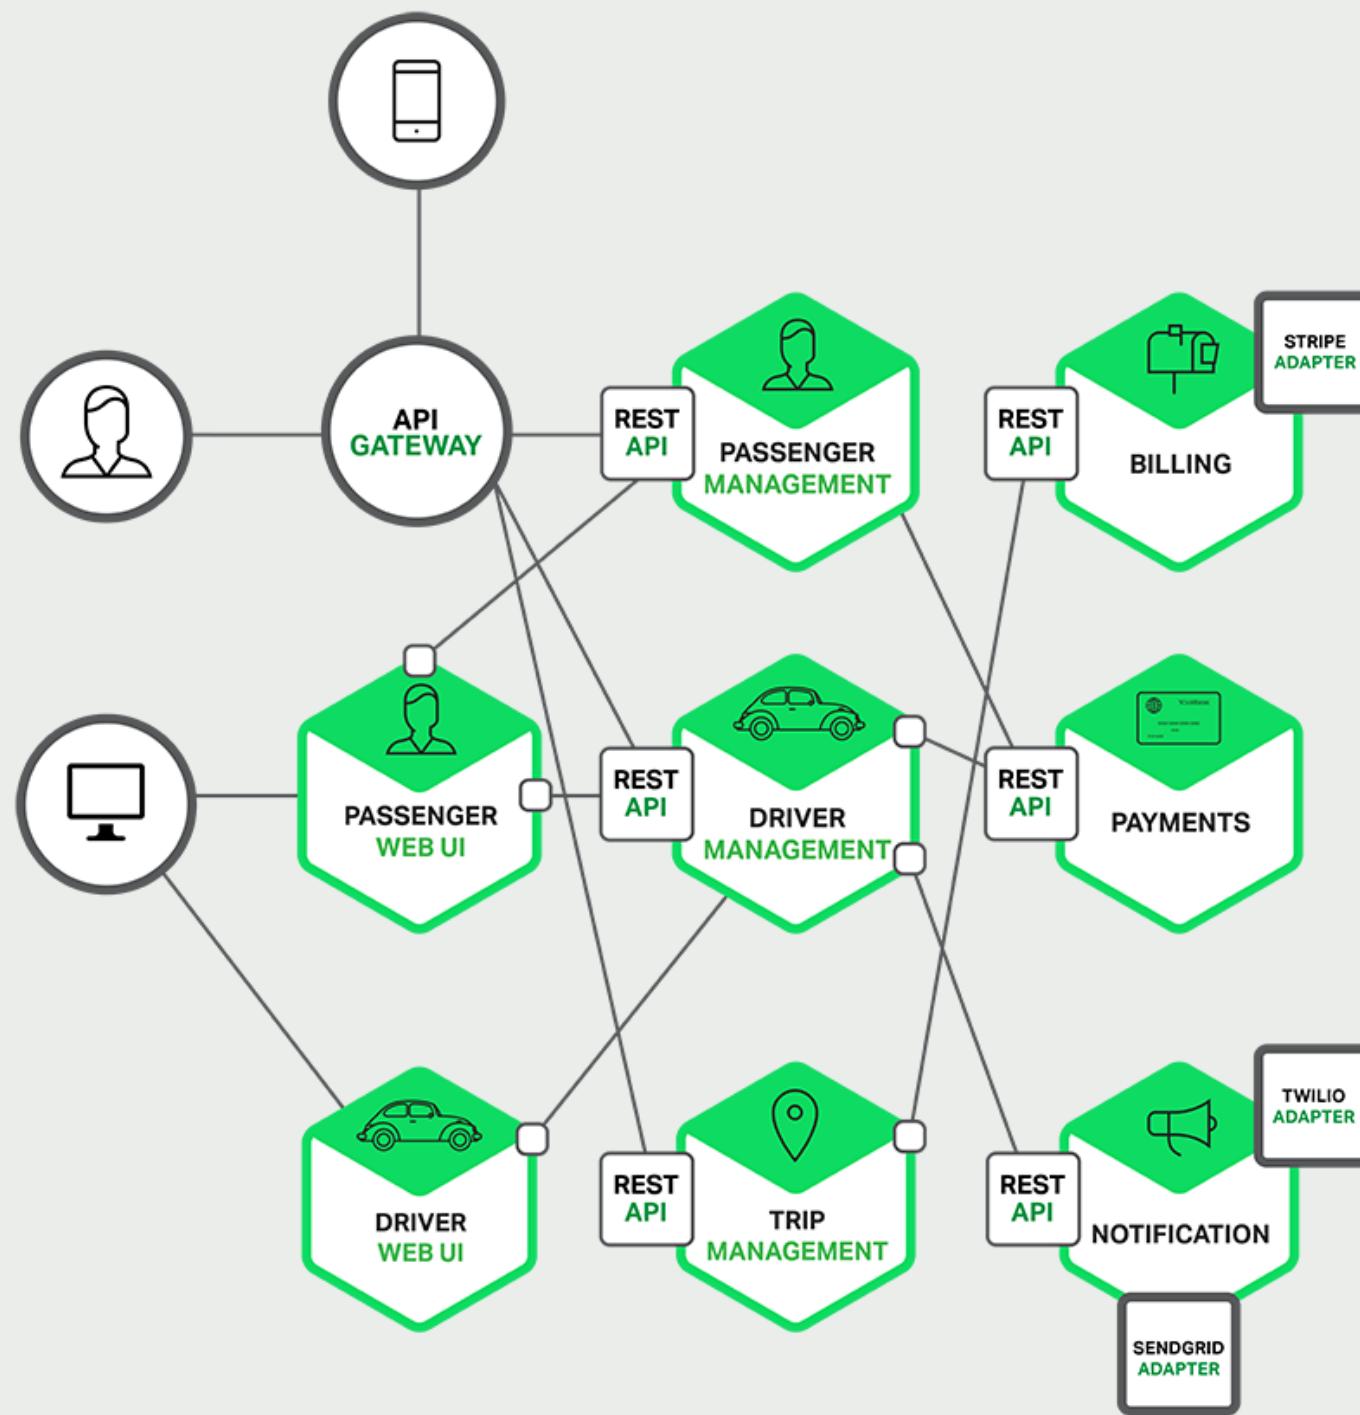
CI/CD

# IDE preparation

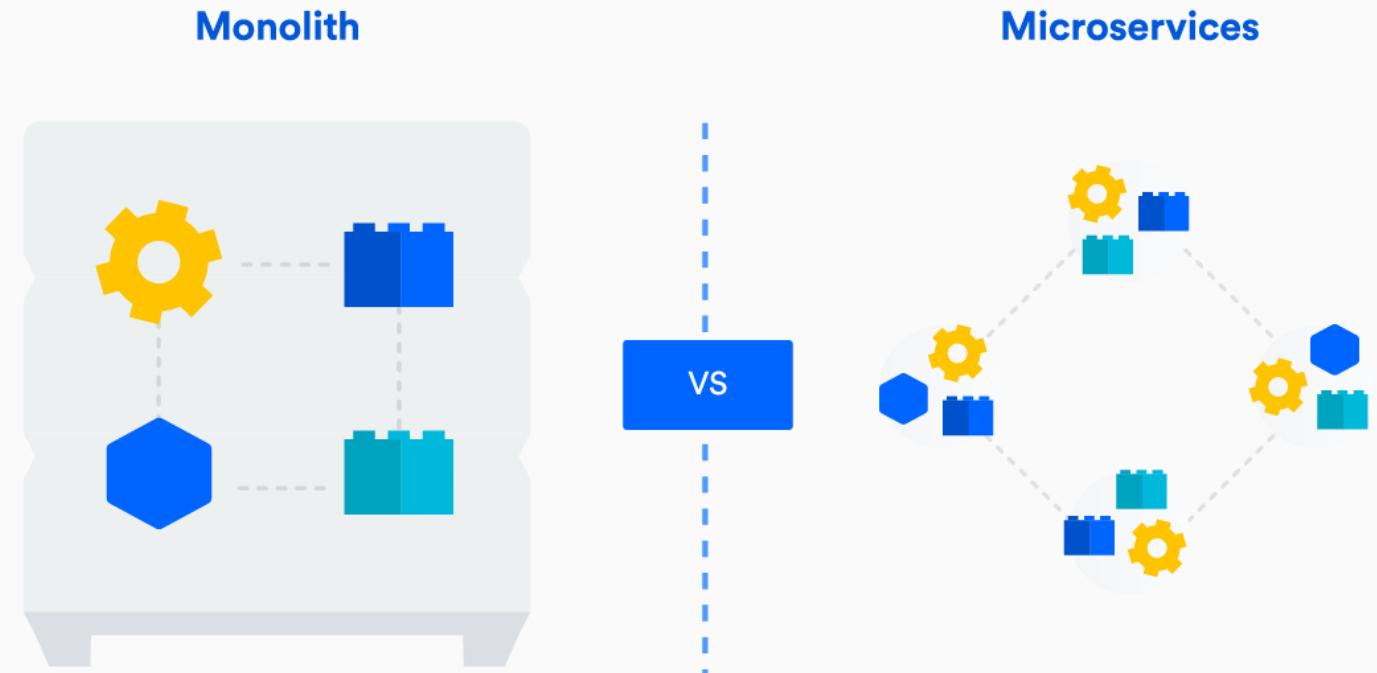
Kilka przydatnych pluginów do VSC:

- Docker
- Kubernetes
- Markdown All in One
- YAML

# Microservices



# Microservice vs Monolith



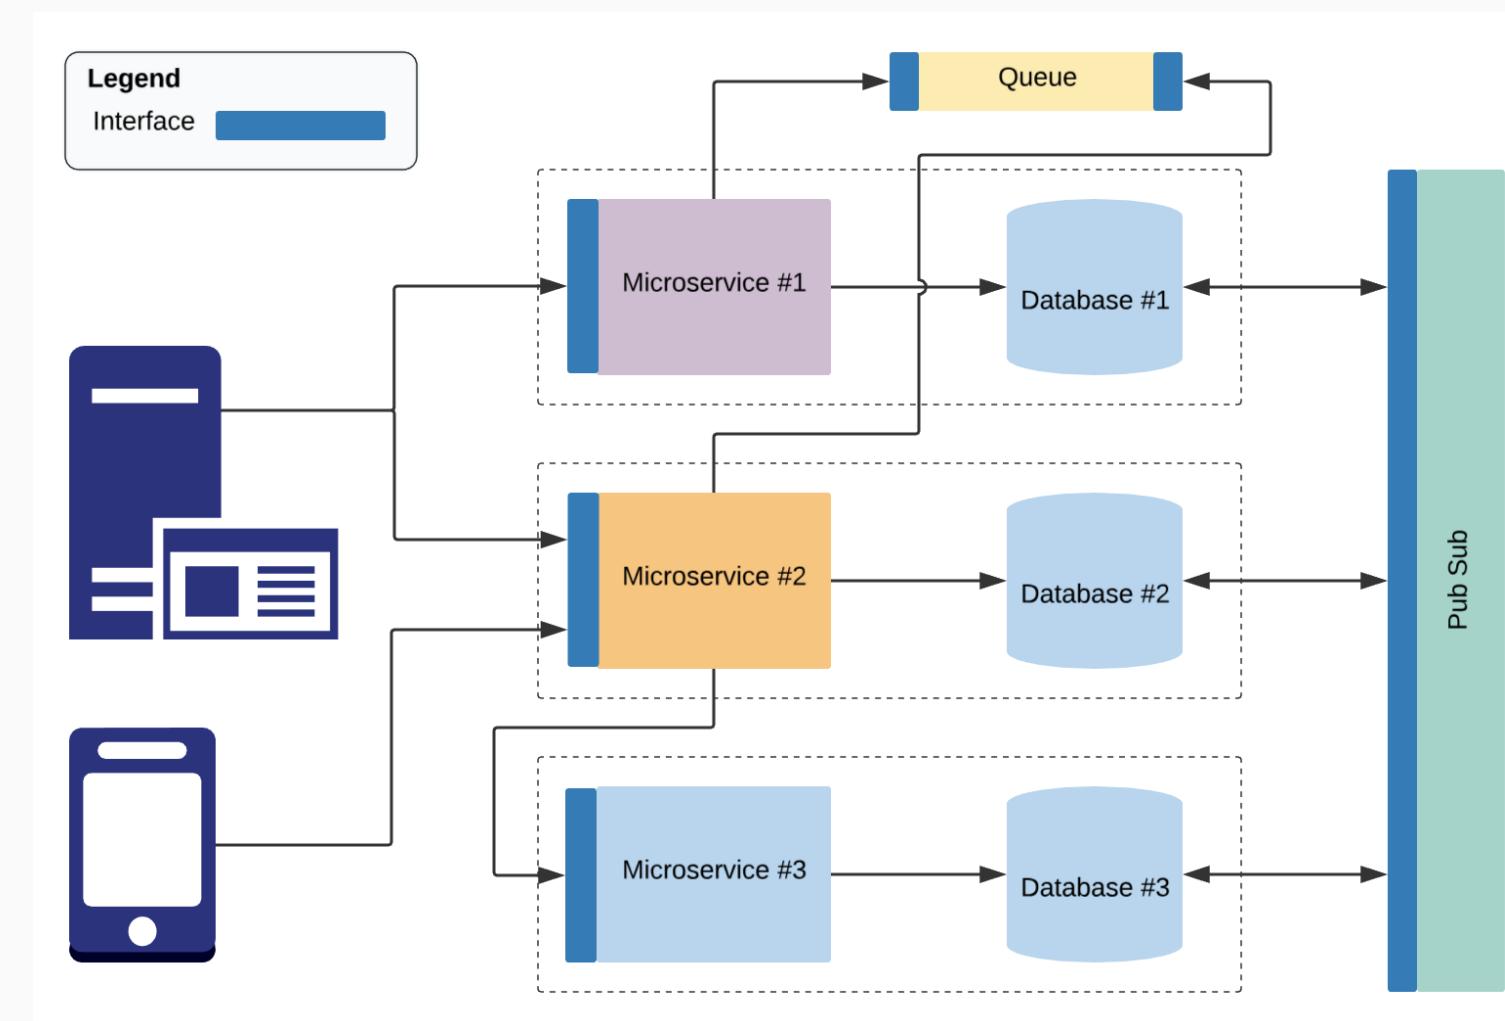
# Microservice as an app

- ls | grep | sed
- Microservice and domain

# Microservice as an app – how to communicate?



# Microservice as an app



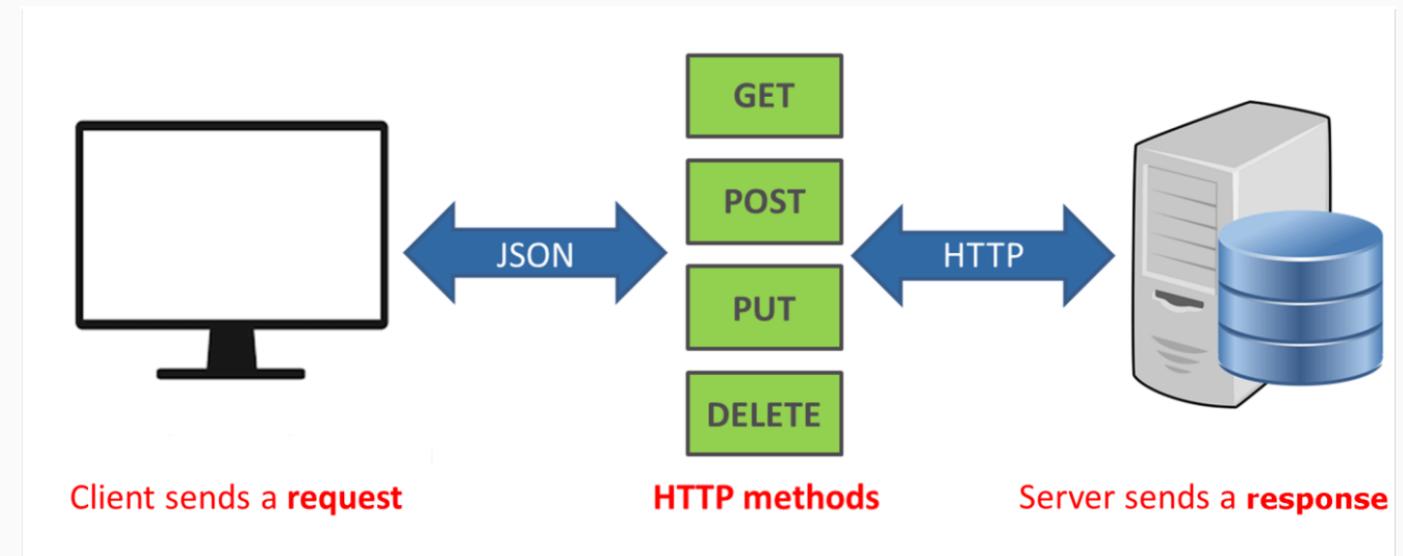
# Microservice as an app

-  Using containers (e.g; Docker) doesn't mean you are adopting Microservices
-  If different services, share code directly within the project, then you are doing it wrong.
-  Depending on too many Microservices to perform a single task, could potentially lead to dependency and performance issues.
-  There is a common misconception using middleware (API Gateway's, Messaging Queues & etc.) Microservices is a bad practice.
-  For certain applications, it is not mandatory for Microservices to communicate via REST APIs.
-  It is also important to understand that when designing Microservices, some of the best practices we used to follow in developing Monoliths could be challenger (eg. Common database).
-  Dividing Microservices based on technical layers also considered as a bad practice.

# Microservice

## – REST API

REST = Representation State Transfer



# REST API Basis



Uniform Interface



Endpoints



Stateless (all information  
must be send)

# REST API

## Basis



Use nouns not verbs (Products vs GetProducts)



Use query for filter  
(Products?name='ABC')



Proper HTTP codes (200, 201, 202, 400, 204/404)



Use versions (v1/products)

# **REST API**

## **Basis**

**GET** – access data

**POST** – insert new data

**PUT** – modify data

**DELETE** – remove data

**PATCH** – partial update

# REST API

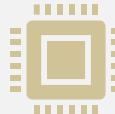
## Maturity



Level 0: Define one URI, and all operations are POST requests to this URI.



Level 1: Create separate URIs for individual resources.



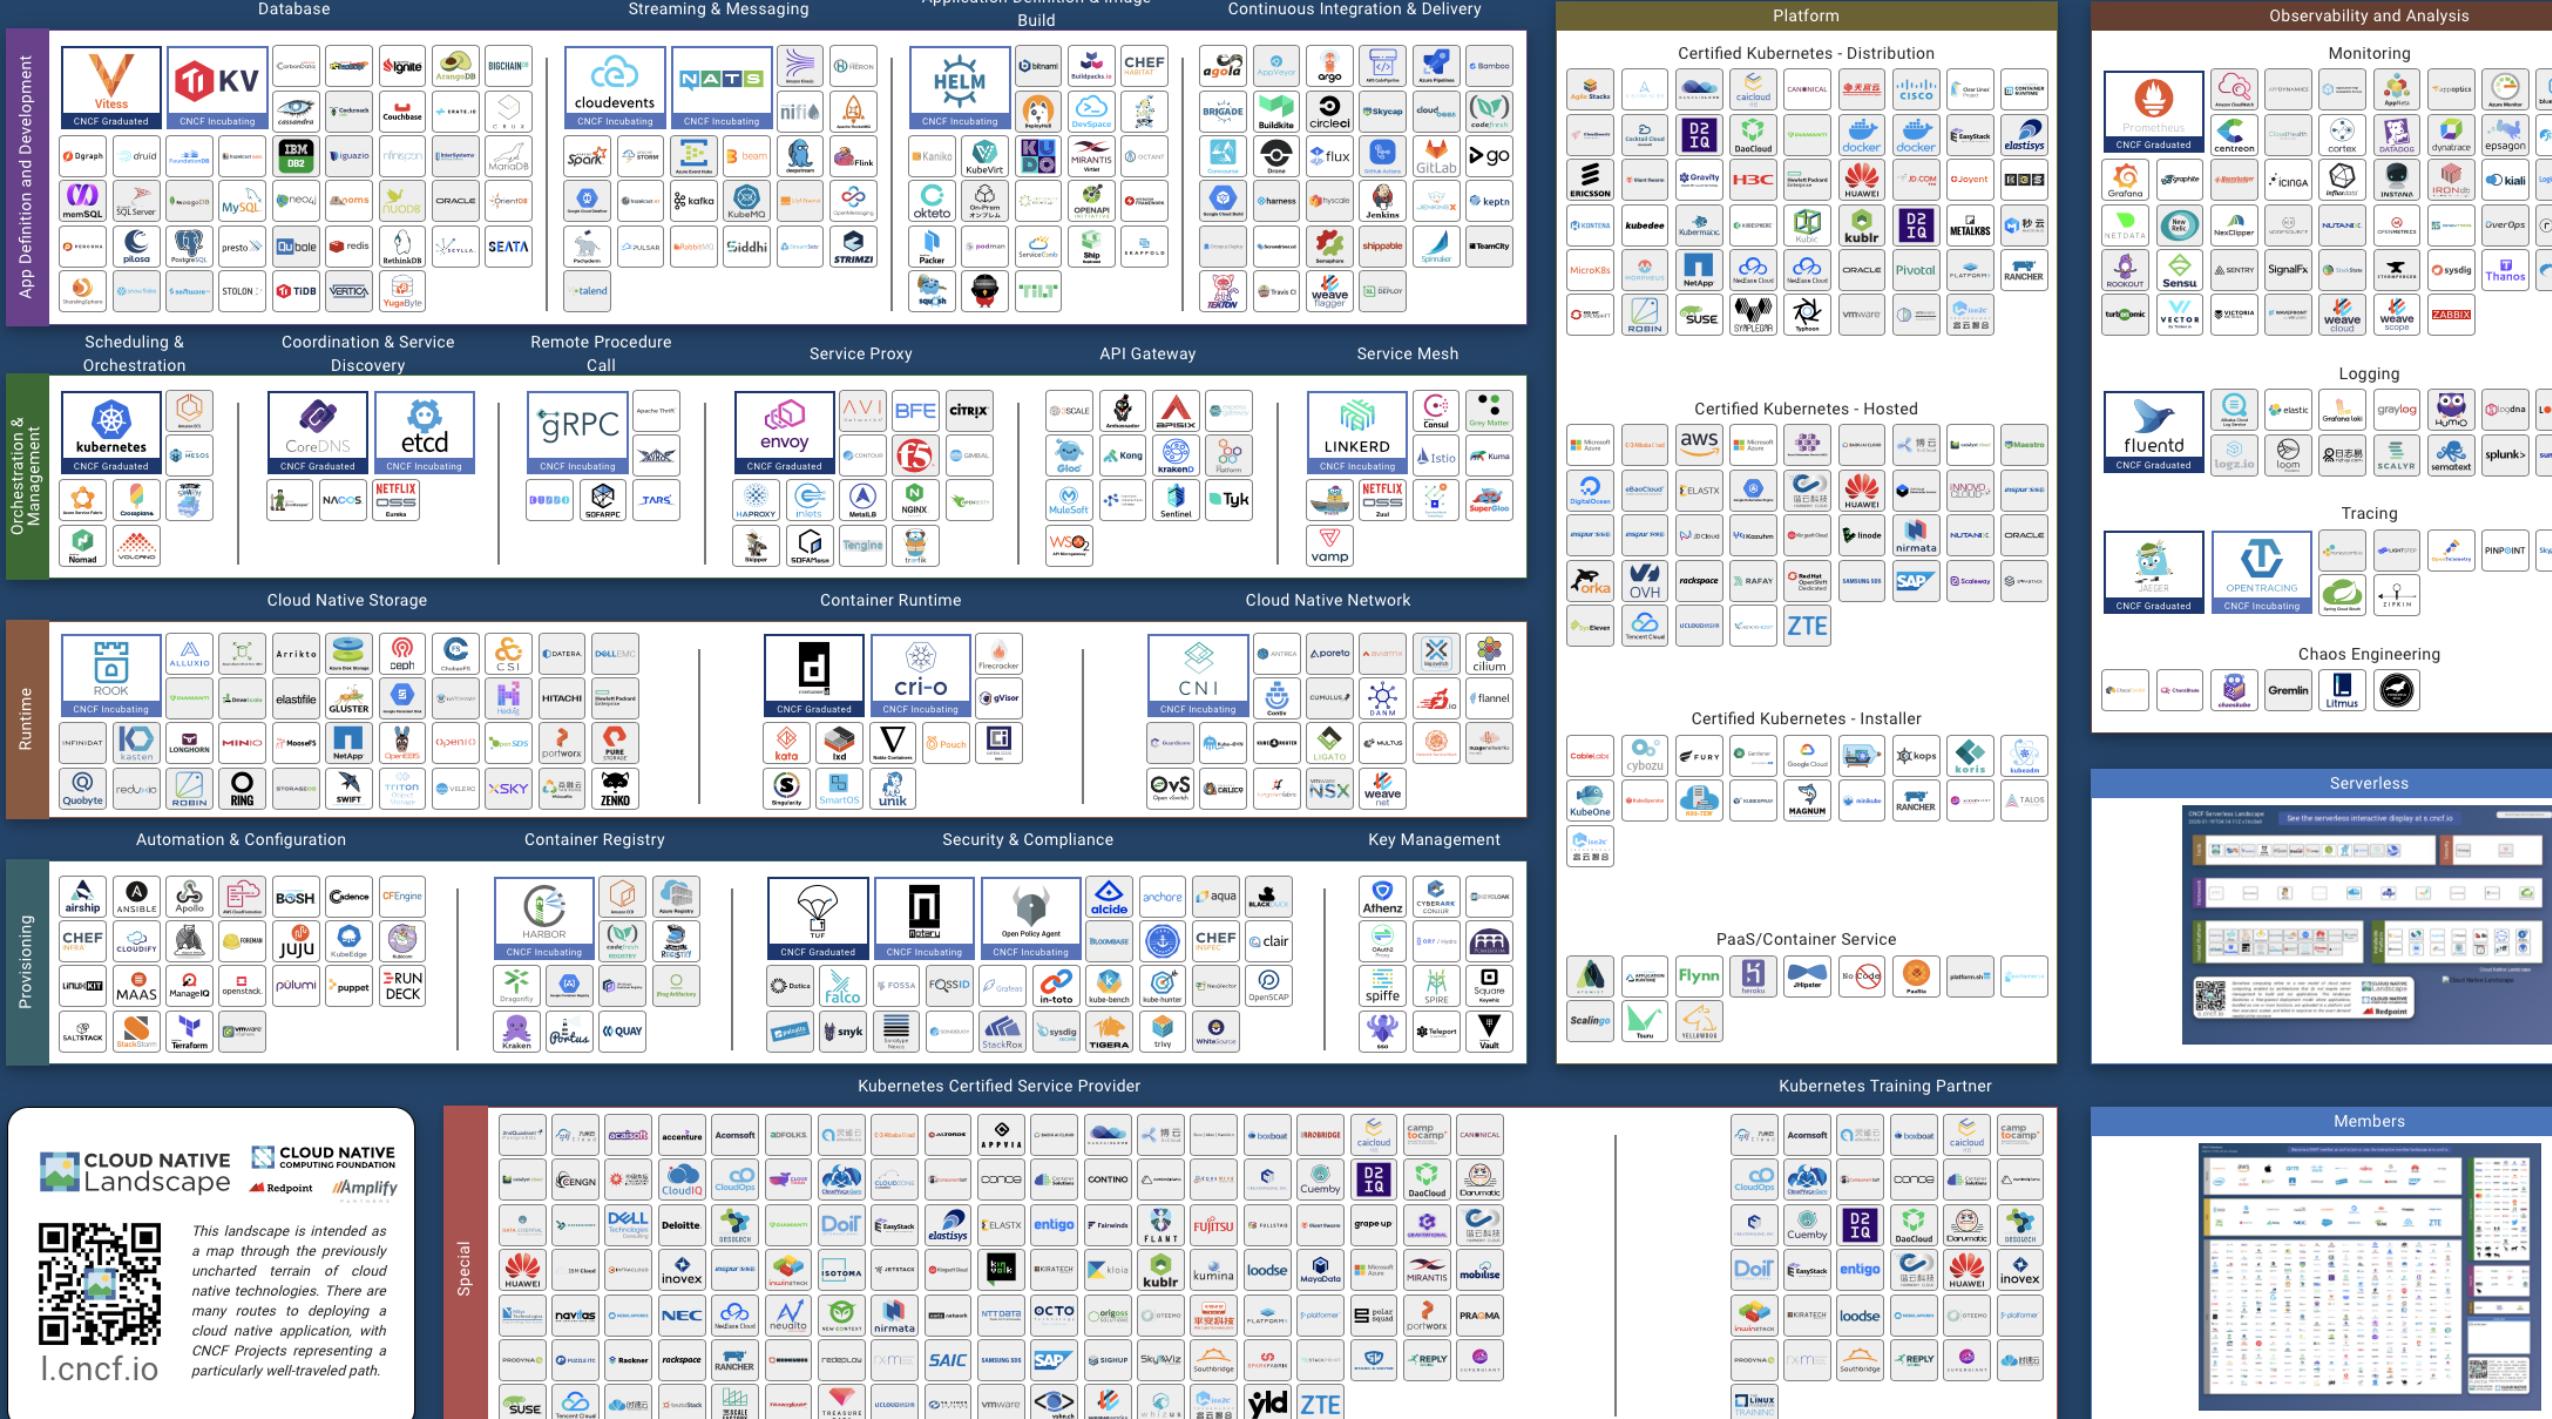
Level 2: Use HTTP methods to define operations on resources.

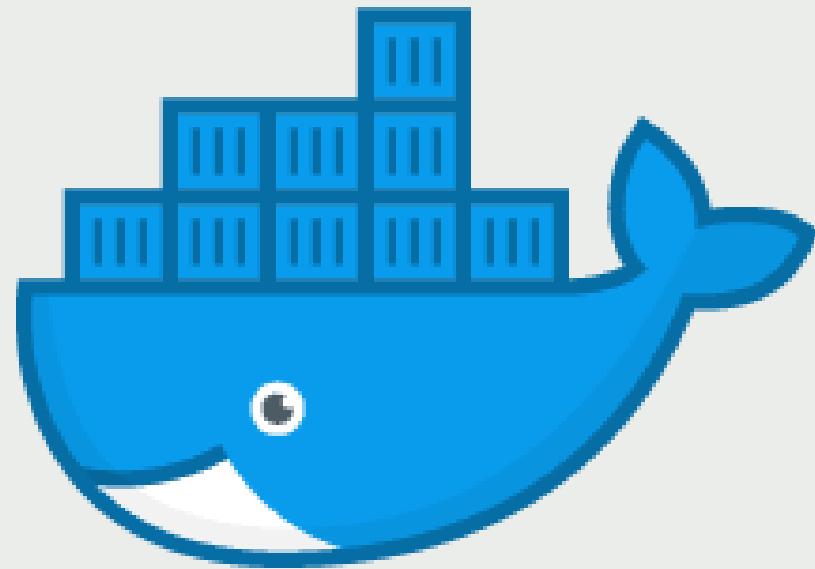


Level 3: Use hypermedia (HATEOAS).

# SWAGGER

<https://petstore.swagger.io>



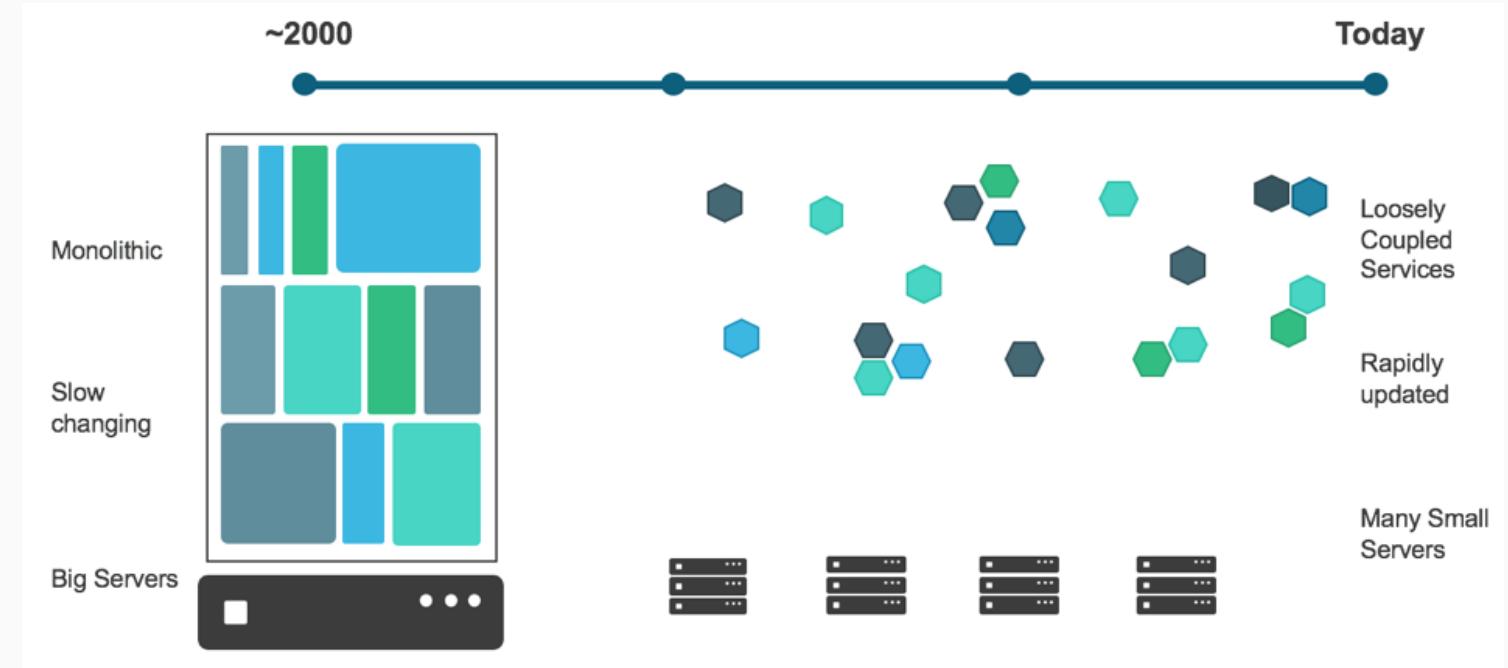


docker

Why  
Docker?



# Why Docker?



# Solution to problems



Different libraries version



Different OS versions



Reproduction of behavior – DEV != STAGE != PROD



Coexistence of software & applications (backward/forward compatibility)



No need for new VM each time new software is installed



Great for SAAS, PAAS solutions



Run Tests same way against DEV, STAGE, PROD (unless you can't)

# Docker – technologies under the hood

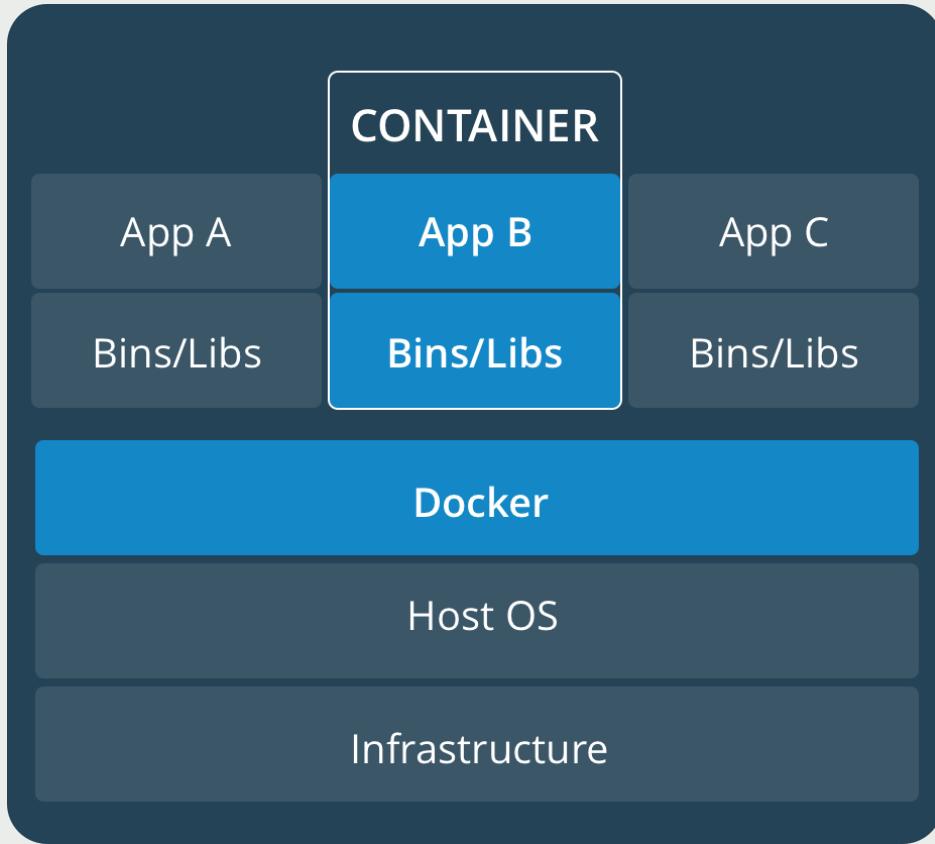
## **namespaces:**

- **The pid namespace:** Process isolation (PID: Process ID).
- **The net namespace:** Managing network interfaces (NET: Networking).
- **The ipc namespace:** Managing access to IPC resources (IPC: InterProcess Communication).
- **The mnt namespace:** Managing filesystem mount points (MNT: Mount).
- **The uts namespace:** Isolating kernel and version identifiers. (UTS: Unix Timesharing System).

**control groups (cgroups)**

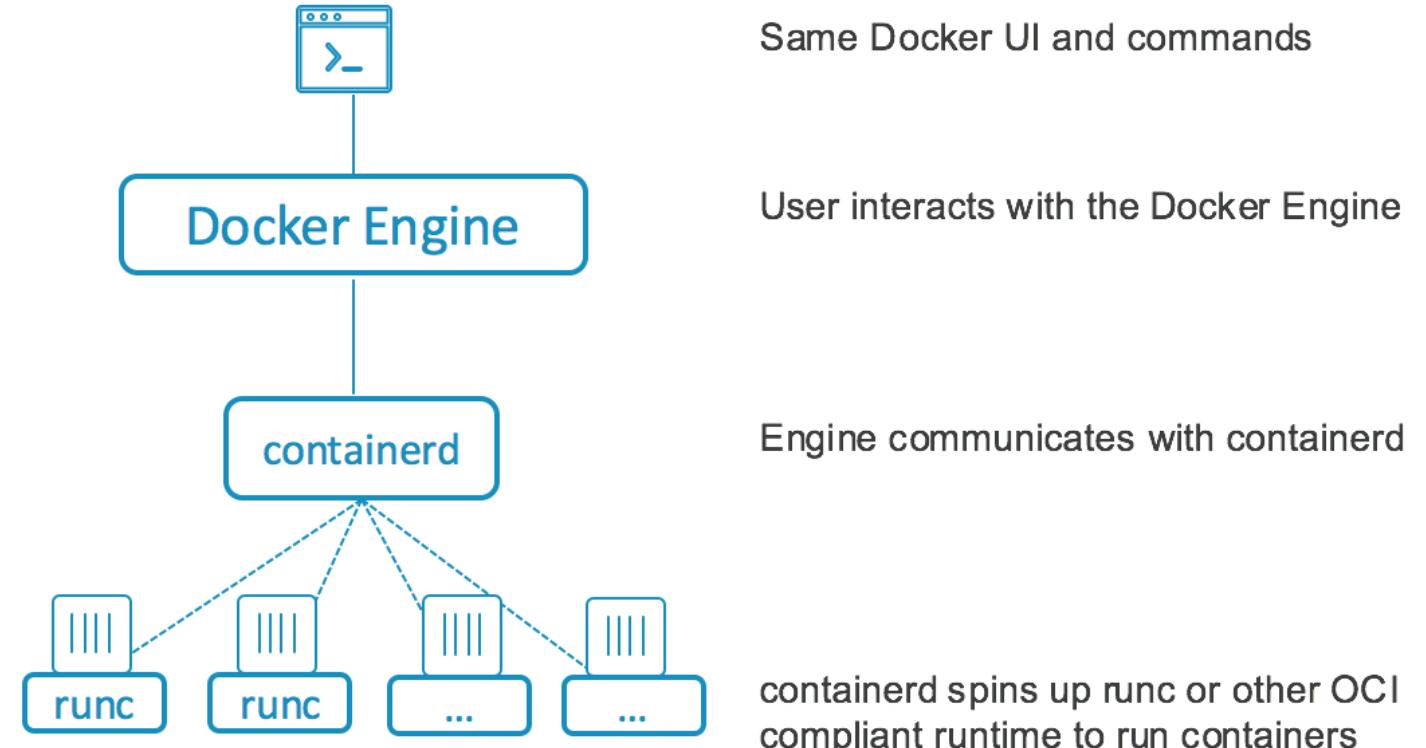
**union file system (AUFS, brtfs, vfs,  
DeviceMapper, overlay)**

**container format (libcontainer)**



# Is docker a VM?

# Docker, containerd?

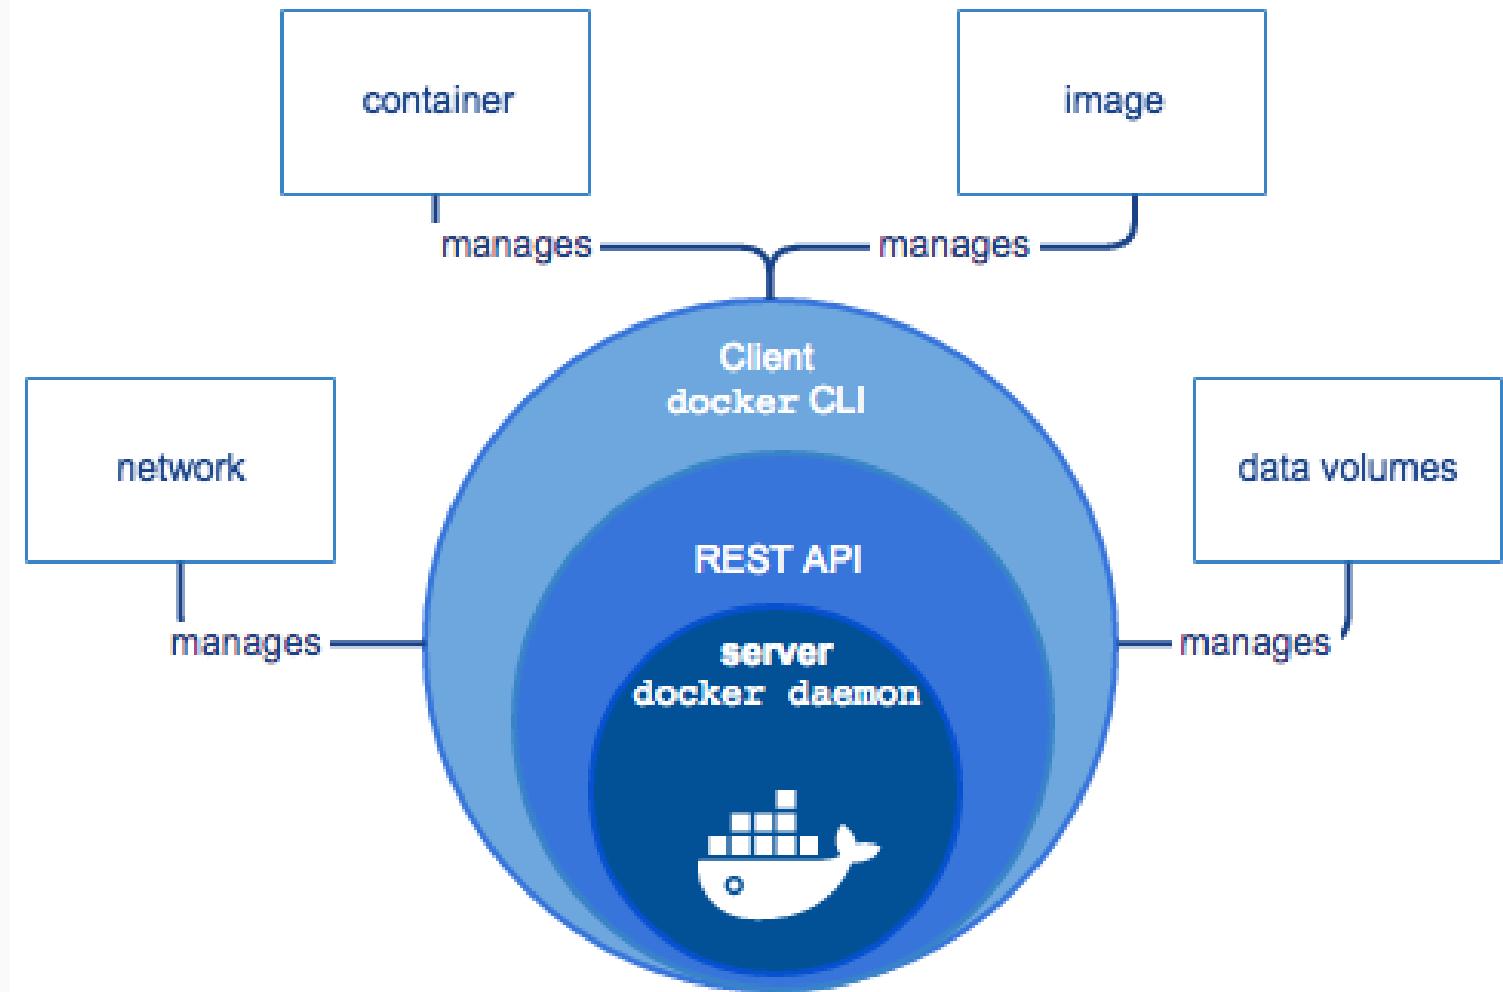


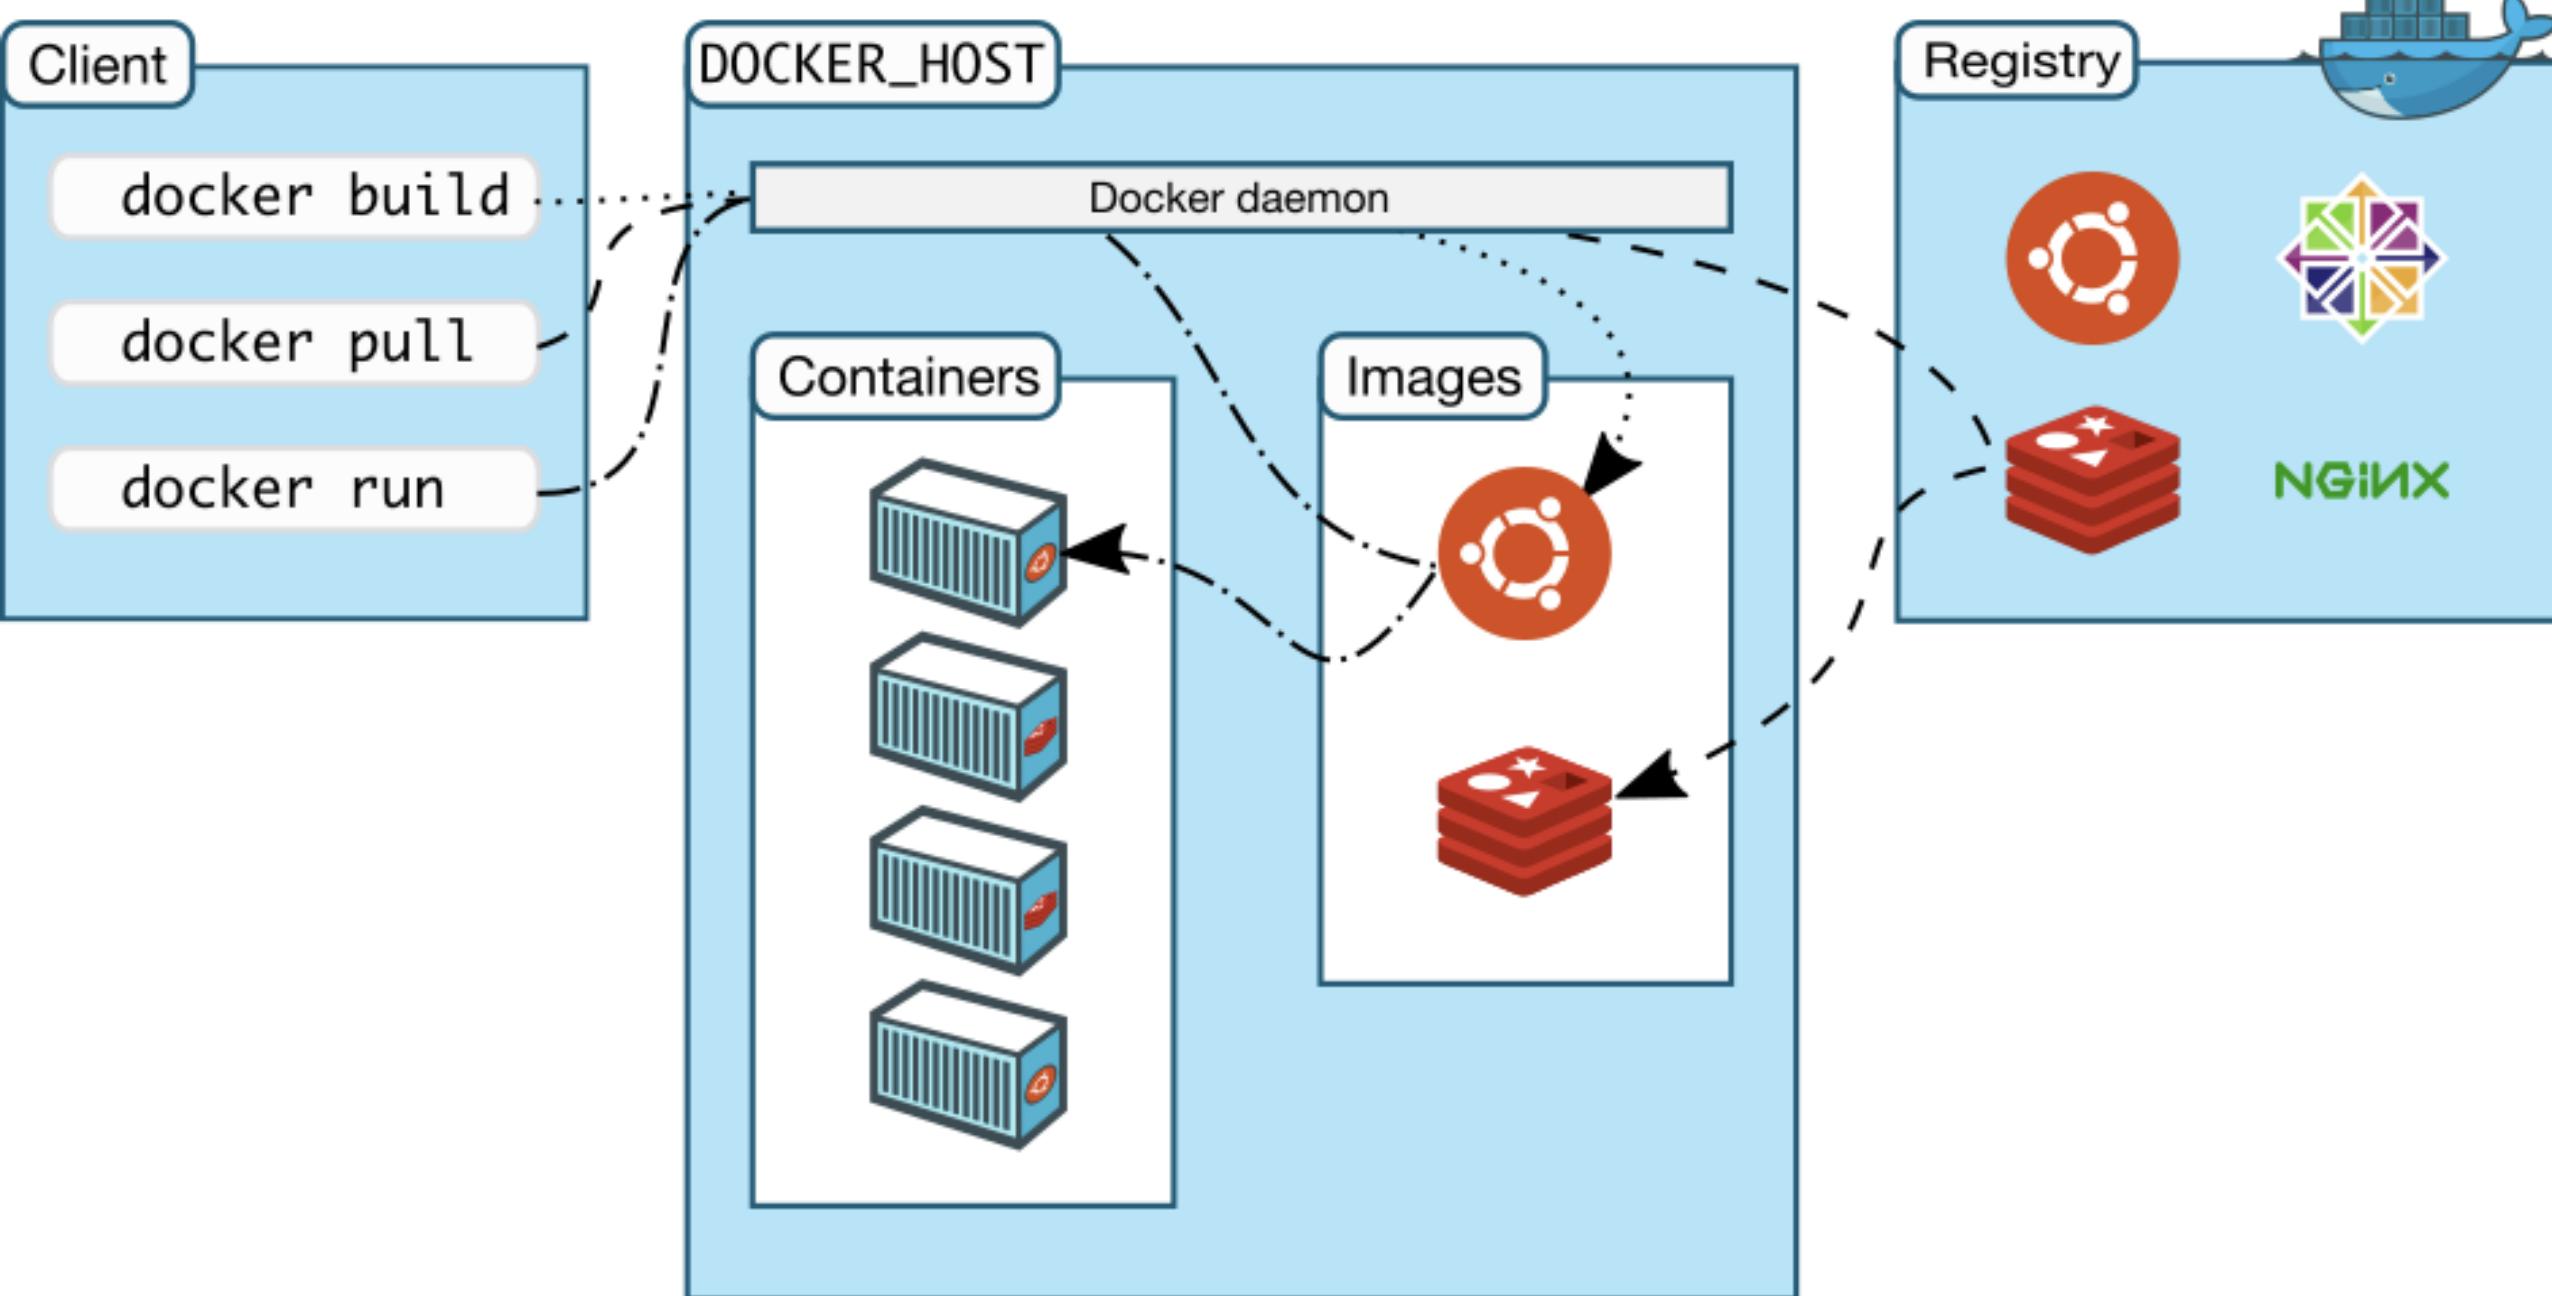
# Docker – the only solution?

OCI Standard:

- runc
  - runtime-spec
  - image-spec
  - distribution-spec
- Containerd
- Podman
- Kata Containers
- gVisor
- CRI-O

# What Docker is?





# Documentation

<http://docs.docker.com>

# Docker - installation

**Wspiera wszystkie OS'y:**

- **Linux**
- **OS X\* (Hyperkit)**
- **Windows\*\* (Hyper-V/WSL)**

Windows Server >= 2016

Windows 10 64-bit: Pro,  
Enterprise, or Education (Build  
17134 or later) or Home\*\*\*.

# Docker - installation

Dla Linux'a

- <https://get.docker.com/>
- <https://docs.docker.com/compose/completion/>

# Docker CLI

```
docker help  
docker <COMMAND> --help
```

# Docker – First Task

```
docker container run hello-world  
docker run hello-world
```

# Dockerfile

- FROM
- LABEL
- RUN
- COPY and ADD
- EXPOSE
- ENV
- ARG
- ENTRYPOINT and CMD
- VOLUME

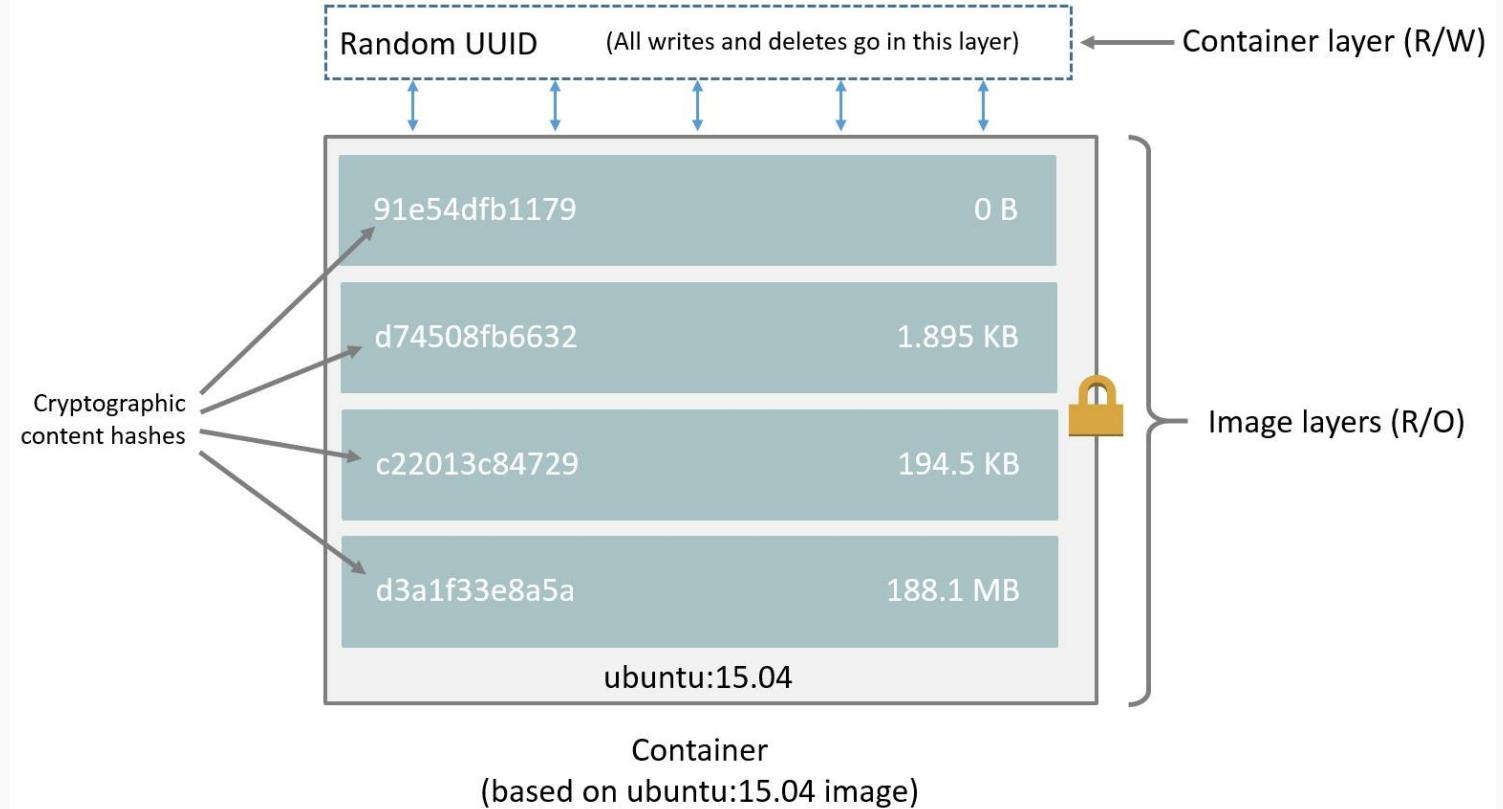
	<b>No ENTRYPOINT</b>	<b>ENTRYPOINT exec_entry p1_entry</b>	<b>ENTRYPOINT ["exec_entry", "p1_entry"]</b>
<b>No CMD</b>	error, not allowed	/bin/sh -c exec_entry p1_entry	exec_entry p1_entry
<b>CMD ["exec_cmd", "p1_cmd"]</b>	exec_cmd p1_cmd	/bin/sh -c exec_entry p1_entry	exec_entry p1_entry exec_cmd p1_cmd
<b>CMD ["p1_cmd", "p2_cmd"]</b>	p1_cmd p2_cmd	/bin/sh -c exec_entry p1_entry	exec_entry p1_entry p1_cmd p2_cmd
<b>CMD exec_cmd p1_cmd</b>	/bin/sh -c exec_cmd p1_cmd	/bin/sh -c exec_entry p1_entry	exec_entry p1_entry /bin/sh -c exec_cmd p1_cmd

# Entrypoint vs CMD

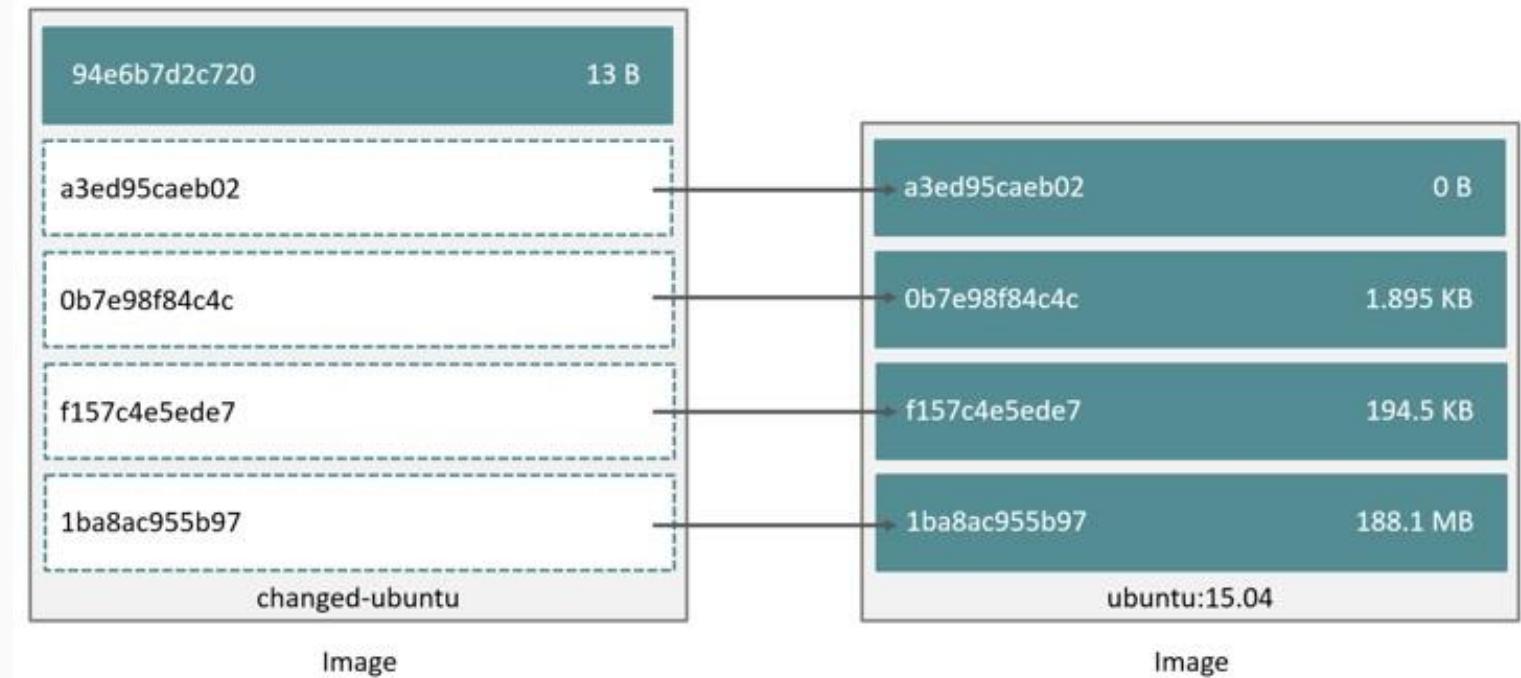
# Docker – .dockerignore

Rule	Behavior
# comment	Ignored.
*/temp*	Exclude files and directories whose names start with temp in any immediate subdirectory of the root. For example, the plain file /somedir/temporary.txt is excluded, as is the directory /somedir/temp.
/*/*temp*	Exclude files and directories starting with temp from any subdirectory that is two levels below the root. For example, /somedir/subdir/temporary.txt is excluded.
temp?	Exclude files and directories in the root directory whose names are a one-character extension of temp. For example, /tempa and /tempb are excluded.

# Docker - Layers



# Docker - Layers



# Docker – Layers verification

```
docker run --rm -it -v \  
/var/run/docker.sock:/var/run/docker.sock \  
wagoodman/dive obraz
```

# Docker – Dockerfile optymalization

<https://github.com/goodwithtech/dockle>

```
docker run --rm -v \  
/var/run/docker.sock:/var/run/docker.sock \  
goodwithtech/dockle wordpress
```

# Docker – Security scan

<https://github.com/aquasecurity/trivy>

```
docker run --rm -v \
/var/run/docker.sock:/var/run/docker.sock \
aquasec/trivy image IMAGE
```

<https://docs.docker.com/engine/scan/>

# Docker – Healthcheck

```
HEALTHCHECK [OPTIONS] CMD command  
HEALTHCHECK --interval=5m --timeout=3s CMD curl -f http://localhost/ || exit 1
```

- interval=DURATION (default: 30s)
- timeout=DURATION (default: 30s)
- start-period=DURATION (default: 0s)
- retries=N (default: 3)

## Exit CODES

- 0: success - the container is healthy and ready for use
- 1: unhealthy - the container is not working correctly
- 2: reserved - do not use this exit code

# Docker – Container



run



attach



**exec -ti      nginx [/bin/bash | /bin/sh]**



stop



start



restart



rm

# Docker – Container



PAUSE == SIGSTOP



STOP == SIGTERM +  
SIGKILL

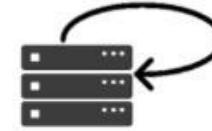


KILL == SIGKILL

1

### Containerize Monoliths

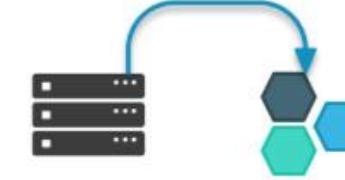
Build-Test for CI; Migrate to the Cloud;  
Achieve better CapEx/OpEx than with VMs



2

### Containerize Monolith; Transform to Microservices

Look for Shared Services to Transform



3

### Enable New Microservices and Apps

Greenfield CaaS



# Docker how to migrate

# Docker – Dockerfile for DB

- easy way:
- Np. Tak jak w obrazie [mariadb](#) poprzez skopiowanie plików do katalogu „/docker-entrypoint-initdb.d”

# Docker – Dockerfile for DB

hard way example #1 (mysql):

```
RUN /bin/bash -c "/usr/sbin/mysqld --bootstrap \
--verbose=0 < file.sql" && rm -f file.sql
```

# Docker – Dockerfile for DB

hard way example #1 (postgres):

```
RUN setuser postgres postgres --single -jE ${POSTGRES_DB} \  
< file.sql; rm -f file.sql
```

# Docker – registry

<https://distribution.github.io/distribution/>

TL;DR

```
docker run -d -p 5000:5000 --name registry registry:2
```

# Docker – registry

```
docker tag ubuntu localhost:5000/myfirstimage  
docker push localhost:5000/myfirstimage
```

<https://distribution.github.io/distribution/about/configuration/>

<https://www.docker.com/pricing>

<https://distribution.github.io/distribution/recipes/mirror/>

Docker lokalne  
repozytorium czy w  
chmurze?

**ONE CONTAINER IS NOT ENOUGH**



**WE NEED TO GO DEEPER**

```
docker run --privileged -it -d docker:stable-dind
```

```
docker run -v /var/run/docker.sock:/var/run/docker.sock -it docker
```

---

DiD

# Docker - ReCap



Image build | rm | prune | --help



Container | create | run | stop... | prune | --help  
Volume | create | rm | prune | --help



Tag image/repository:version  
dst\_address/repository:version | --help



Build it -> Run it -> Tag it -> Push it

# Docker – bigger number of containers

Docker natively has 2 tools that support the management of bigger number of containers:

- compose
- swarm

**Docker –  
bigger  
number of  
containers**

Docker compose – configuration  
<https://docs.docker.com/compose/compose-file/>

# Docker – bigger number of containers

Setting up variables:

**\$ {VARIABLE:-default}** will set a variable if VARIABLE is unset or empty

**\$ {VARIABLE-default}** will set the variable if VARIABLE is unset

Required variables:

**\$ {VARIABLE:?Err}** fails if VARIABLE is unset or empty

**\$ {VARIABLE?Err}** fails if VARIABLE is unset

Using YAML Anchors

<https://docs.docker.com/compose/compose-file/#extension-fields>

# Docker – Logging

Using docker we have two options of managing logs:

- Standard way – each application logs to a local file, logs are rotated, etc.
- Logs are redirected to STDOUT

<https://docs.docker.com/engine/admin/logging/overview/>





# Documentation

<http://docs.docker.com>

# Docker – Best practices

Consider one Dockerfile per folder

Use GIT!

Keep image small

Don't put secrets! (or any other 'personal' stuff)

Execute one app per container

Use .dockerignore

Layers are your friend – use them as middle step before build

# Docker – Production Ready

Docker Host only

Avoid Docker0 „cloud”

Consider Haproxy, Nginx

Healthz / Healthchecks

Use ELK / GELF / LOKI stacks for logs

Use clustered services

Consider host discovery / dns automation (Consul)

Consider Network Driver

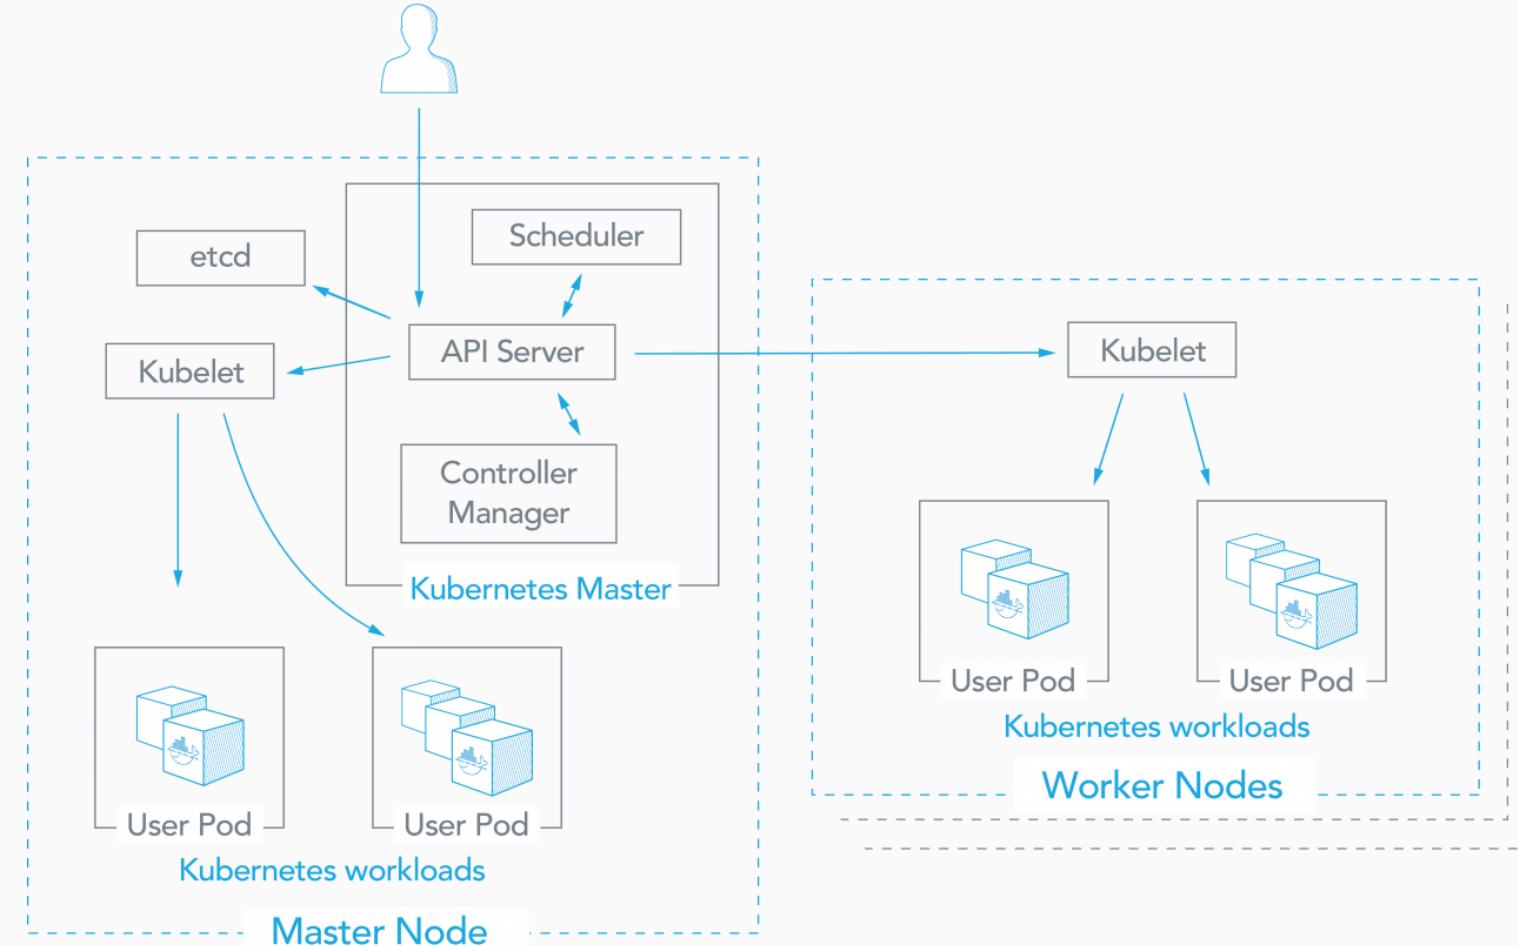
BUT USE K8S!



**kubernetes**

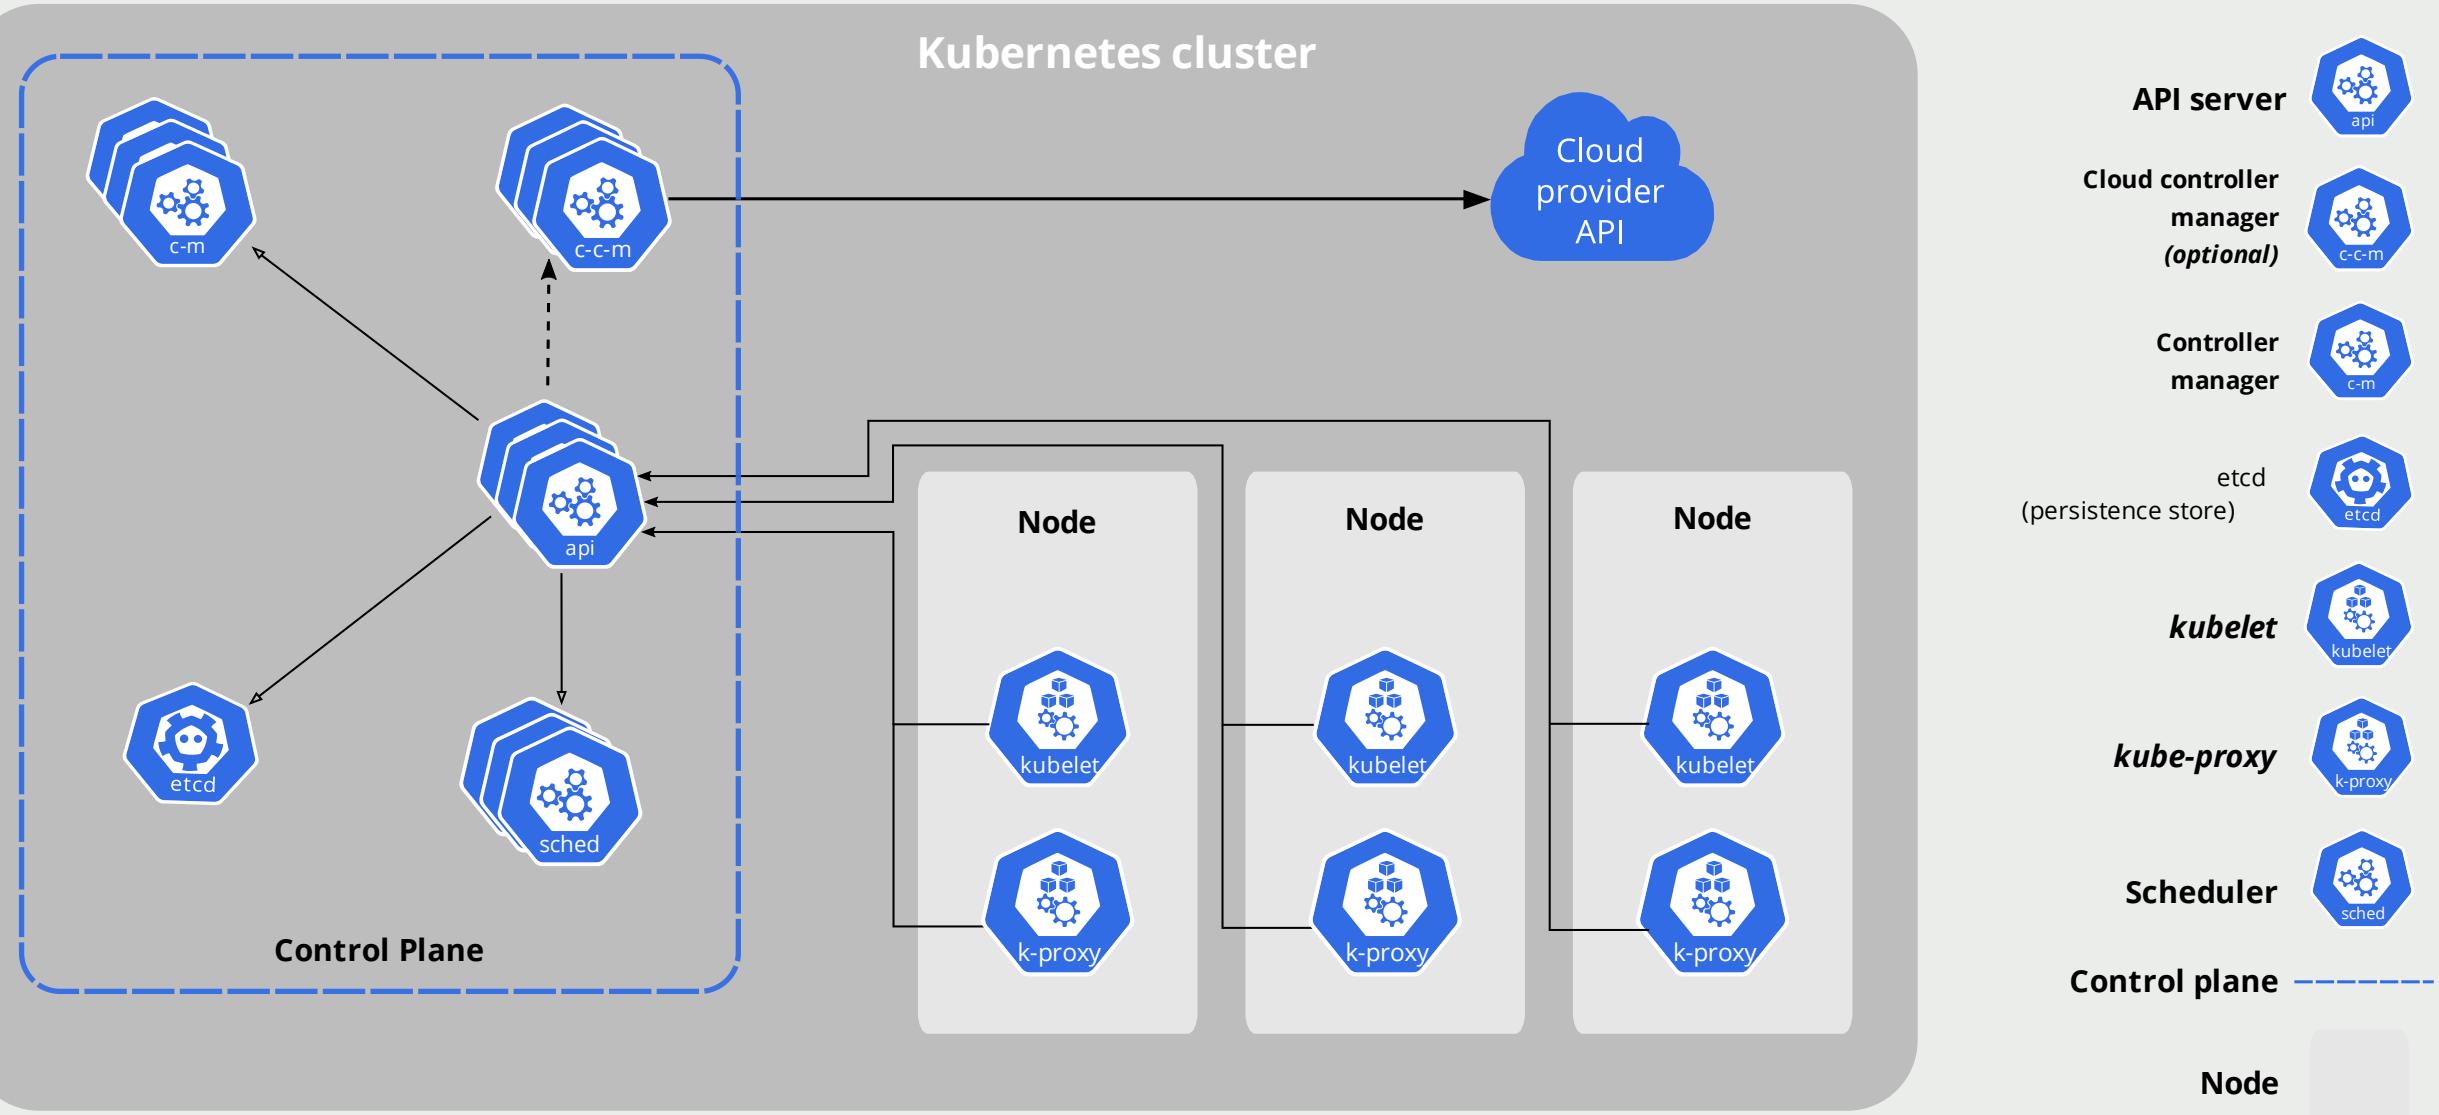
# Kubernetes

## Architecture



API Server: management hub for Kubernetes  
Scheduler: places a workload on the appropriate Node  
Controller Manager: scales workloads up/down  
etcd: stores configuration data which can be accessed by API Server

Kubelet: Receives pod specifications from API Server, updates Nodes  
Master Node: places workloads on Nodes  
Worker Nodes: receives requests from Master Nodes and dispatches them  
User Pod: a group of containers with shared resources



**Kubernetes  
is just a  
tool!**

Building blocks can be  
binaries or docker  
containers or ...

There is no „one setup for  
all”

# Kubernetes

---

## Orchestrator concepts



Replicas



Ingress



Healthcheck



Pod/Services



Configs



Secrets



Volumes

# Kubernetes

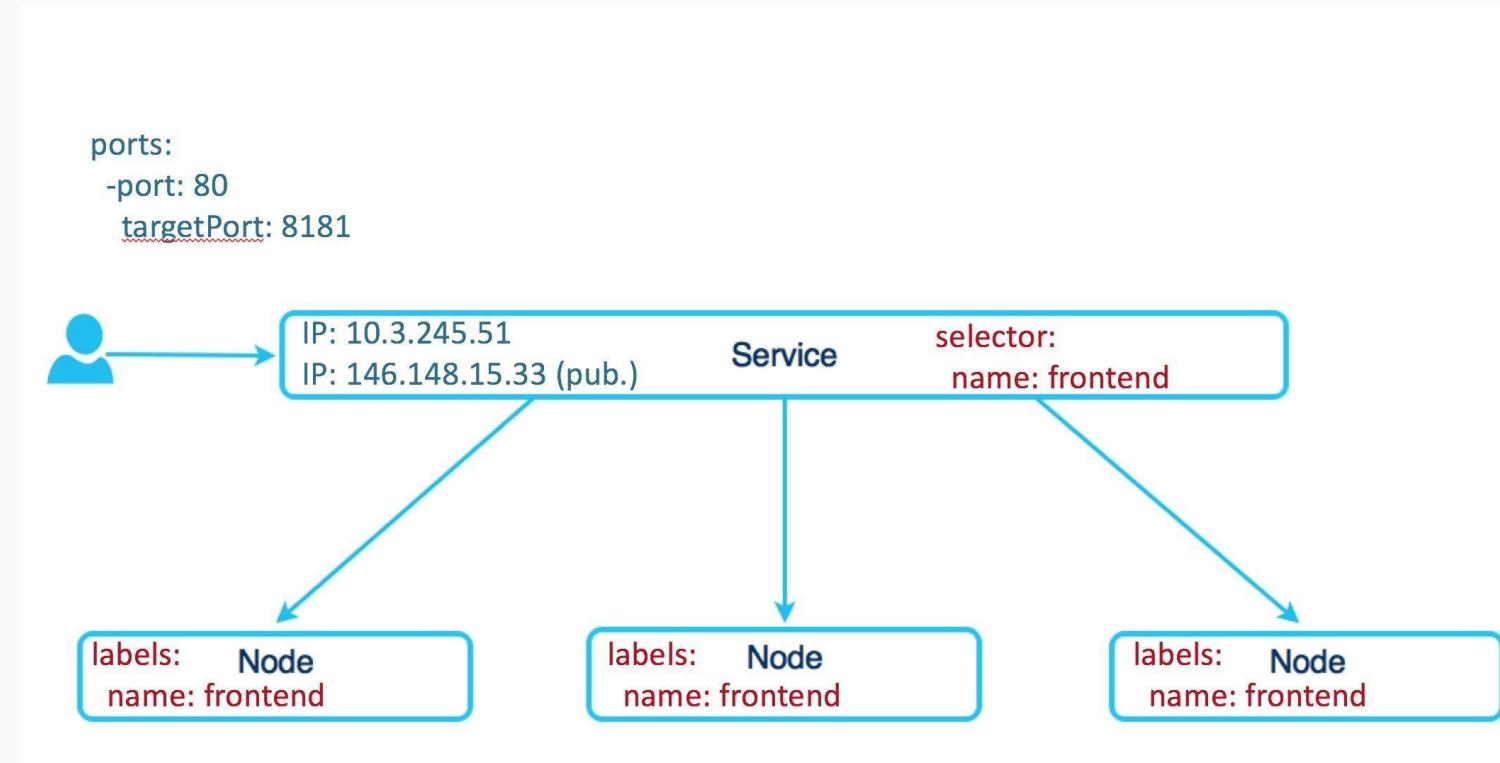
## - kubectl

<https://kubernetes.io/docs/tasks/access-application-cluster/configure-access-multiple-clusters/>

```
$ kubectl config view  
$ kubectl config use-context  
$ kubectl create -f ./<DIR>  
$ kubectl create -f <URL>  
$ kubectl create -R -f ./<DIR>  
$ kubectl port-forward POD PORT:PORT  
$ kubectl exec -ti .. /bin/bash  
$ kubectl logs  
$ kubectl --kubeconfig=./config -n monitoring get pods
```

# Kubernetes

## - Service



# Kubernetes

## - Ingress



# Kubernetes

## – Logic objects



Deployments



StatefulSet (for stateful services like redis)



DaemonSet



Job



CronJob

# Kubernetes

---

## Deployment



Rolling Update



Canary release

# Kubernetes

-

## Deployment



LivenessProbe



ReadinessProbe



Detach from service when  
fail!

# Kubernetes – QoS



Burstable

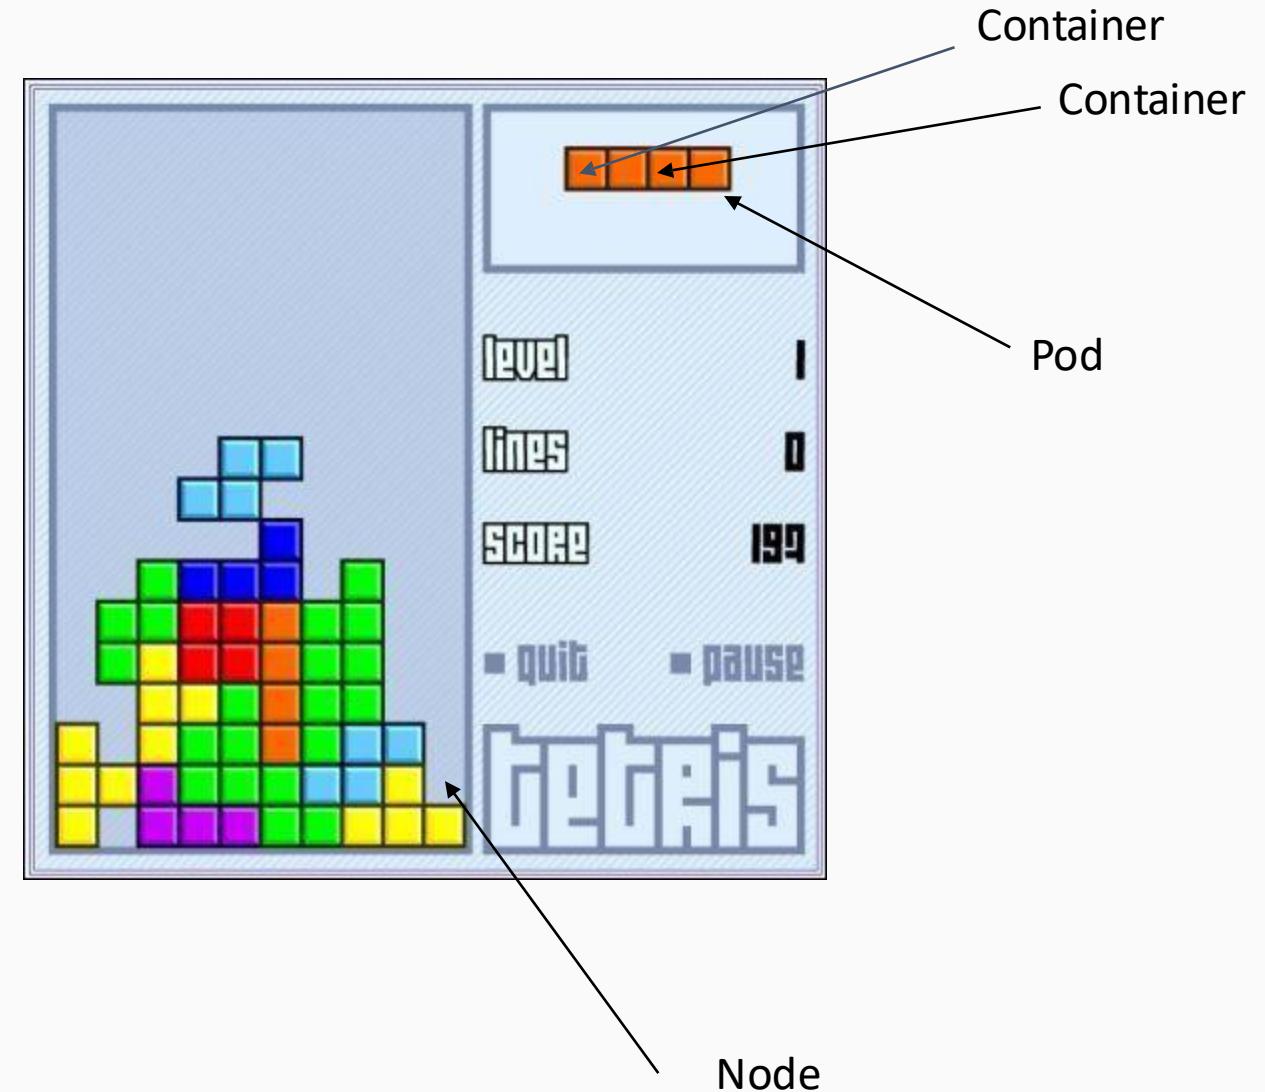


Guaranteed



Best effort

**Kubernetes  
is a player  
trying to  
win!**

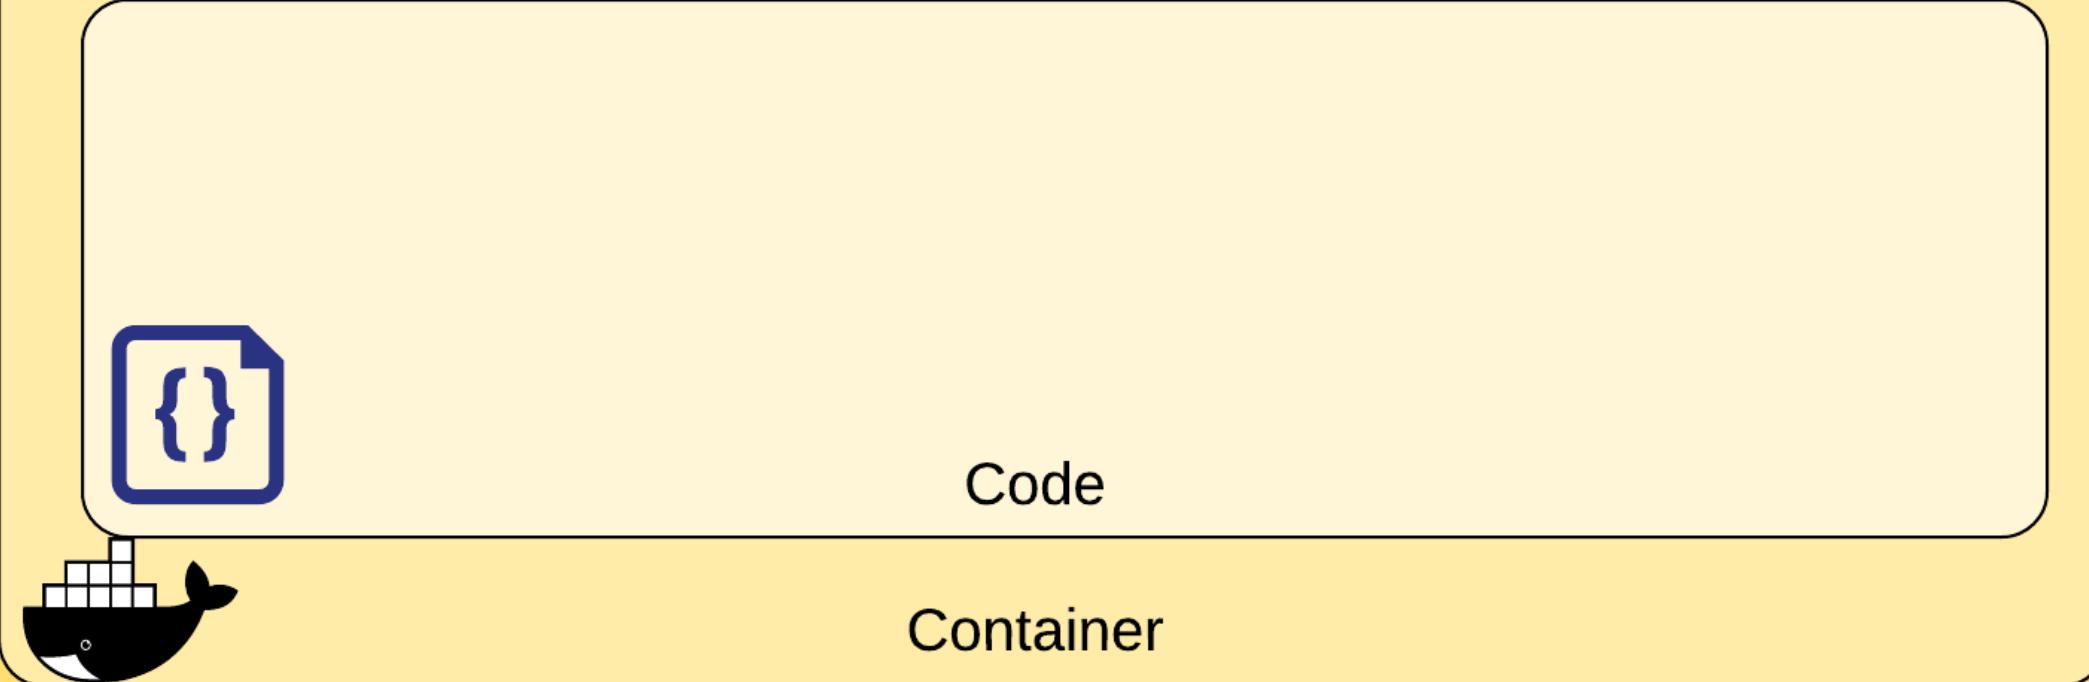


# Kubernetes

---

## Dashboard

<https://kubernetes.io/docs/tasks/access-application-cluster/web-ui-dashboard/>



Computers and  
Networks

Cloud/Co-Lo/Corporate  
Datacenter

# Kubernetes – Cloud Security concerns



Network access to API Server (Control plane)



Network access to Nodes (nodes)



Kubernetes access to Cloud Provider API



Access to etcd



etcd Encryption

# Kubernetes

## – Cluster

## Security

## concerns



RBAC Authorization (Access to the Kubernetes API)



Authentication



Application secrets management (and encrypting them in etcd at rest)



Pod Security Policies



Quality of Service (and Cluster resource management)



Network Policies



TLS For Kubernetes Ingress

# Kubernetes – Container Security concerns



Container Vulnerability Scanning  
and OS Dependency Security



Image Signing and Enforcement



Disallow privileged users

# Kubernetes

## – Code Security concerns



Access over TLS only



Limiting port ranges of communication



3rd Party Dependency Security



Static Code Analysis



Dynamic probing attacks

Initial Access	Execution	Persistence	Privilege Escalation	Defense Evasion	Credential Access	Discovery	Lateral Movement	Collection	Impact
Using Cloud credentials	Exec into container	Backdoor container	Privileged container	Clear container logs	List K8S secrets	Access the K8S API server	Access cloud resources	Images from a private registry	Data Destruction
Compromised images in registry	bash/cmd inside container	Writable hostPath mount	Cluster-admin binding	Delete K8S events	Mount service principal	Access Kubelet API	Container service account		Resource Hijacking
Kubeconfig file	New container	Kubernetes CronJob	hostPath mount	Pod / container name similarity	Access container service account	Network mapping	Cluster internal networking		Denial of service
Application vulnerability	Application exploit (RCE)	Malicious admission controller	Access cloud resources	Connect from Proxy server	Applications credentials in configuration files	Access Kubernetes dashboard	Applications credentials in configuration files		
Exposed Dashboard	SSH server running inside container				Access managed identity credential	Instance Metadata API	Writable volume mounts on the host		
Exposed sensitive interfaces	Sidecar injection				Malicious admission controller		Access Kubernetes dashboard		

### = New technique

= Deprecated technique

# Kubernetes

## – Security concerns



<https://www.cisecurity.org/benchmark/kubernetes/>

CIS Distribution Independent Linux

v2.0.0 - 07-16-2019

<https://github.com/aquasecurity/kube-bench>

# Kubernetes

## – instalacja

<https://kubernetes.io/docs/setup/production-environment/tools/>

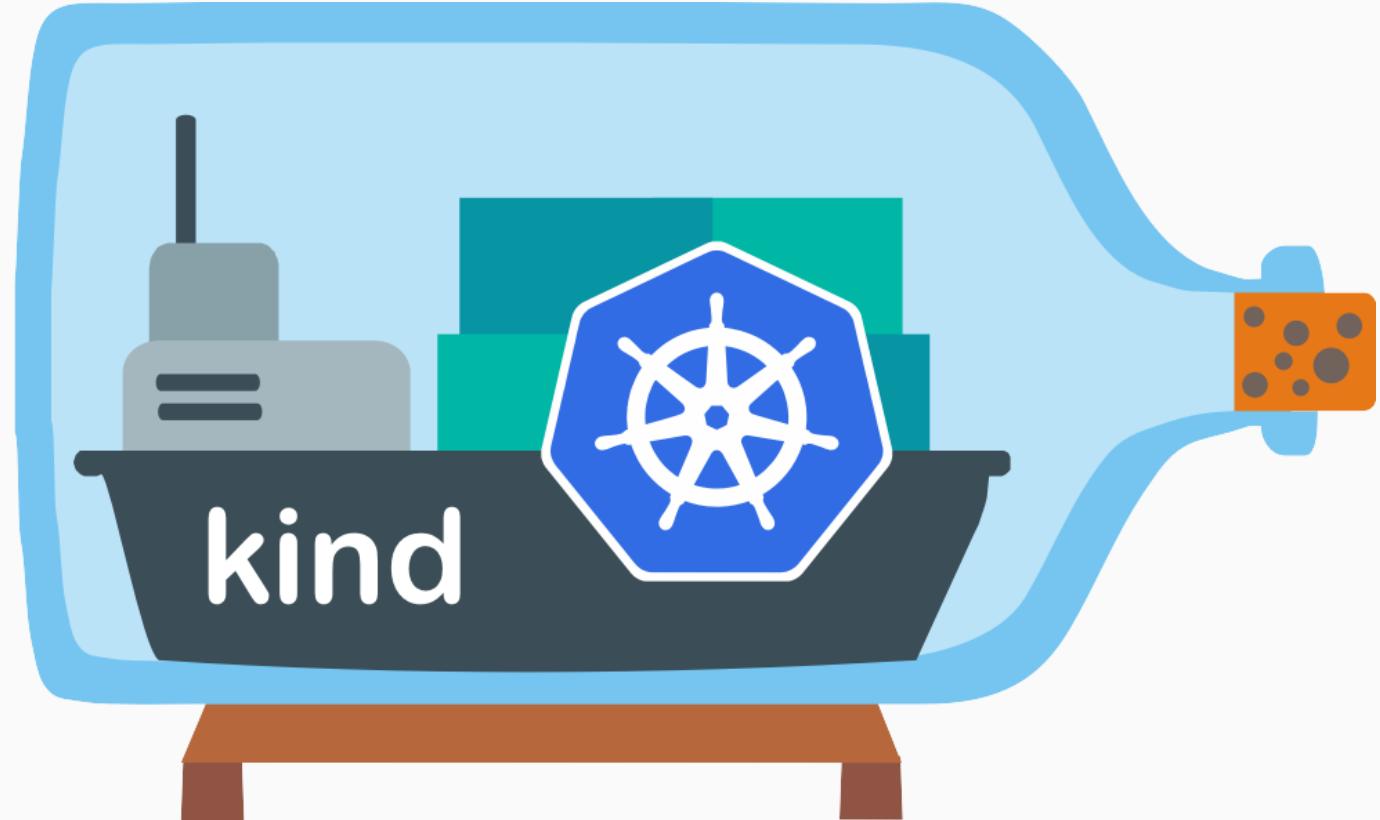
# Kubernetes – instalacja



<https://rancher.com/>

# Kubernetes

– kind



<https://kind.sigs.k8s.io>

# Kubernetes

## – Narzędzia ułatwiające pracę z K8s

<https://k8slens.dev>

<https://www.telepresence.io>

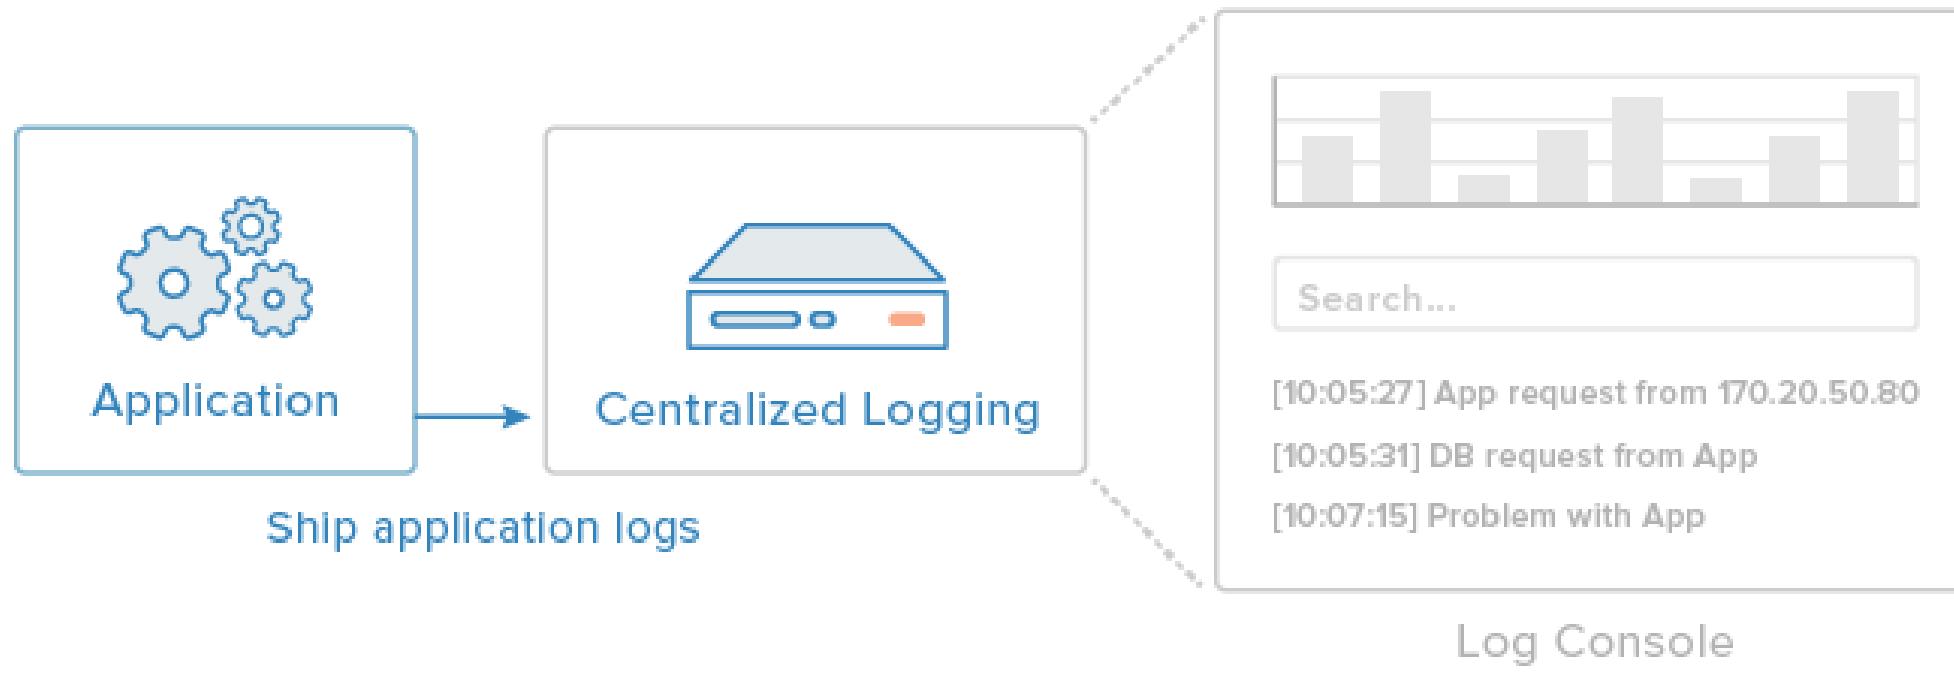
# Documentation

<https://kubernetes.io/docs/home/>

# Documentation

<https://learnk8s.io/production-best-practices>

## Centralized Logging



# Jak przesyłać logi?

- Syslog
- API:
  - Elasticsearch
  - Graylog
  - Loki
  - Rozwiązania chmurowe

# Logowanie – podstawowe pojęcia

- Rotacja
- Retencja

# Rodzaje logów

- Systemowe (OS)
- Aplikacyjne (nginx, postgres itd.)
- Naszych aplikacji (to co napisliśmy)
- Audytowe

Local or centralized?

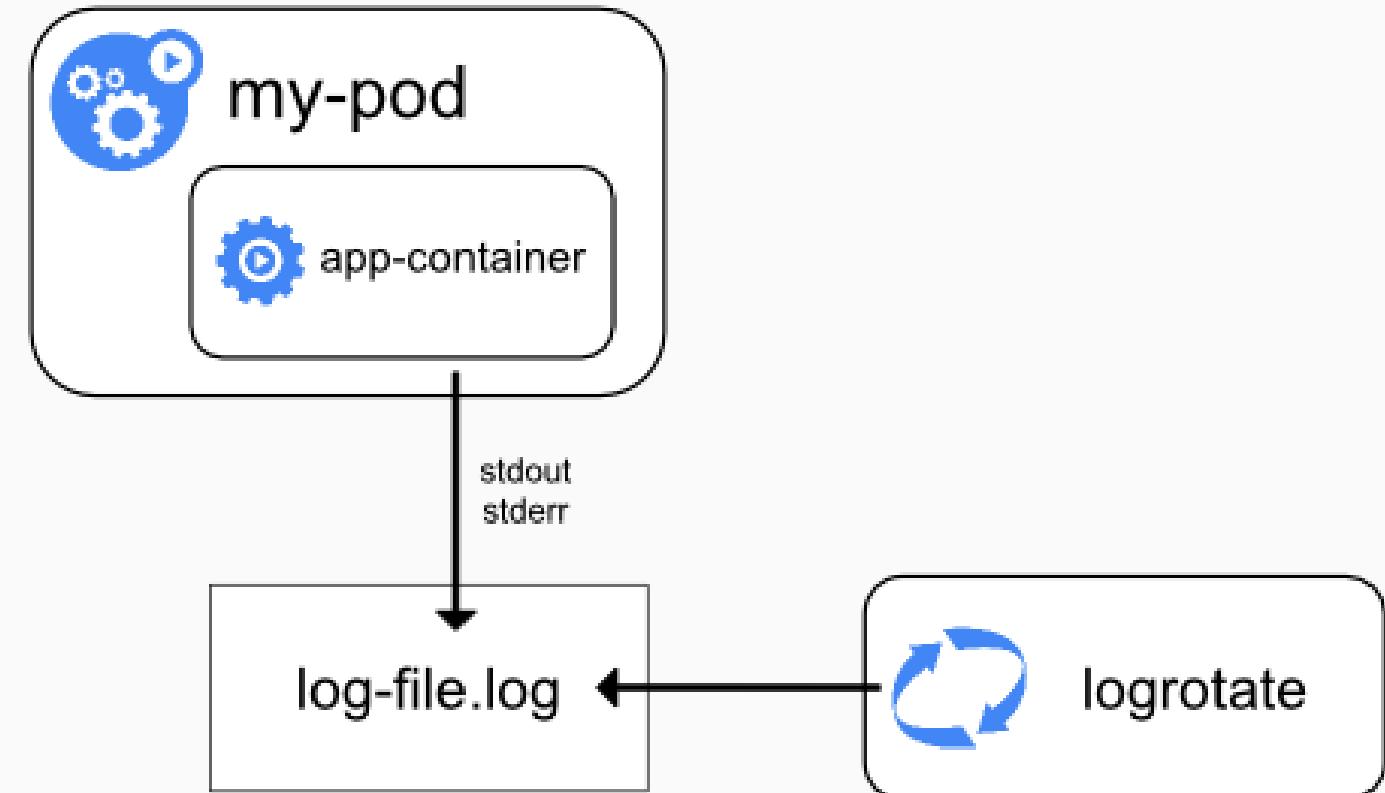
# Docker - Logging

Driver	Description
none	No logs are available for the container and docker logs does not return any output.
<a href="#">local</a>	Logs are stored in a custom format designed for minimal overhead.
<a href="#">json-file</a>	The logs are formatted as JSON. The default logging driver for Docker.
<a href="#">syslog</a>	Writes logging messages to the syslog facility. The syslog daemon must be running on the host machine.
<a href="#">journald</a>	Writes log messages to journald. The journald daemon must be running on the host machine.
<a href="#">gelf</a>	Writes log messages to a Graylog Extended Log Format (GELF) endpoint such as Graylog or Logstash.
<a href="#">fluentd</a>	Writes log messages to fluentd (forward input). The fluentd daemon must be running on the host machine.
<a href="#">awslogs</a>	Writes log messages to Amazon CloudWatch Logs.
<a href="#">splunk</a>	Writes log messages to splunk using the HTTP Event Collector.
<a href="#">etwlogs</a>	Writes log messages as Event Tracing for Windows (ETW) events. Only available on Windows platforms.
<a href="#">gcplogs</a>	Writes log messages to Google Cloud Platform (GCP) Logging.
<a href="#">logentries</a>	Writes log messages to Rapid7 Logentries.

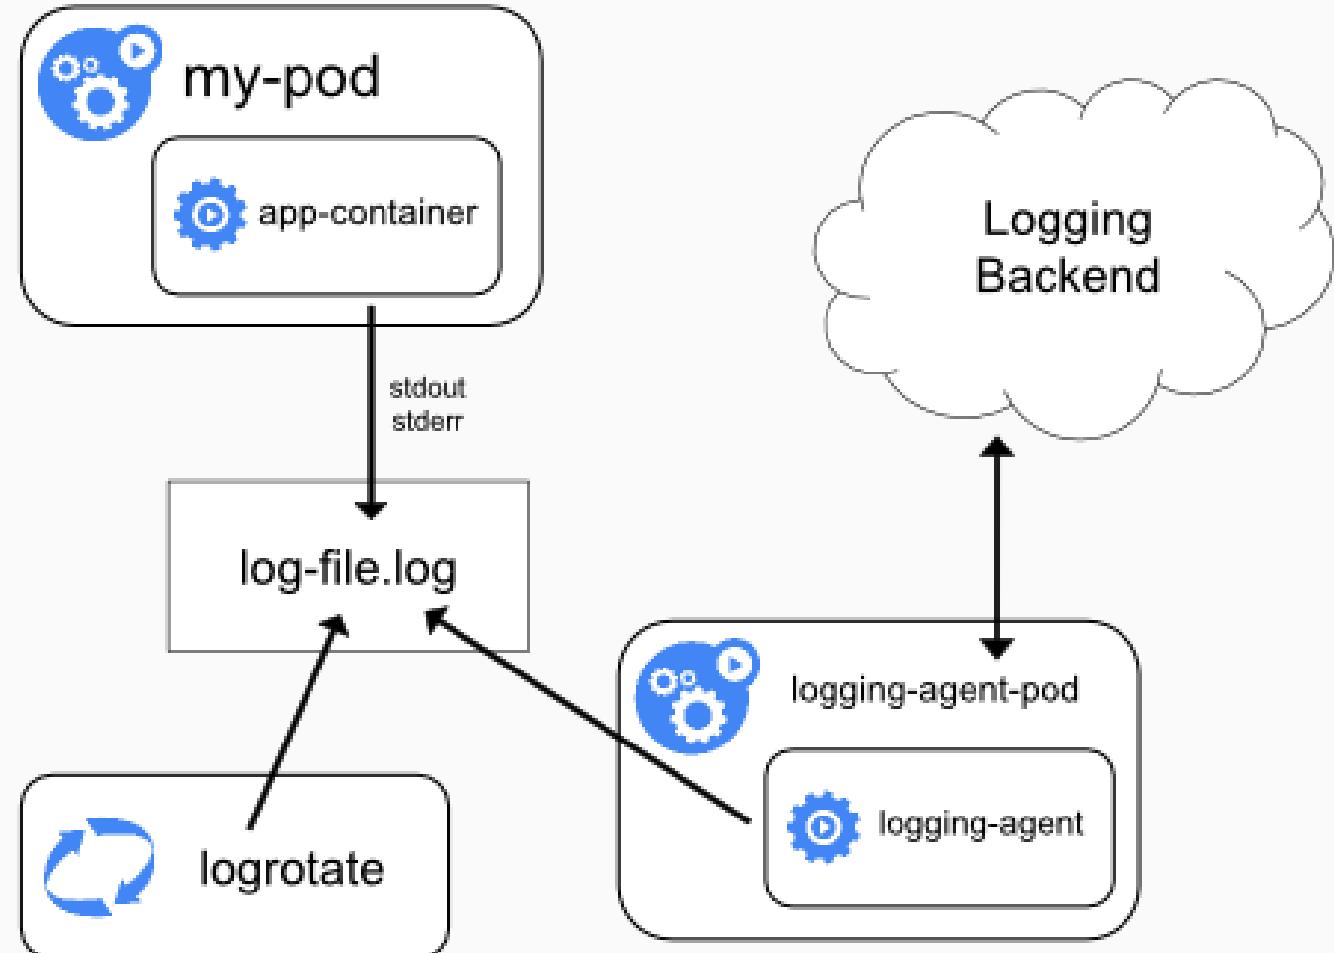
<https://docs.docker.com/config/containers/logging/configure/>

# Kubernetes

- Logging

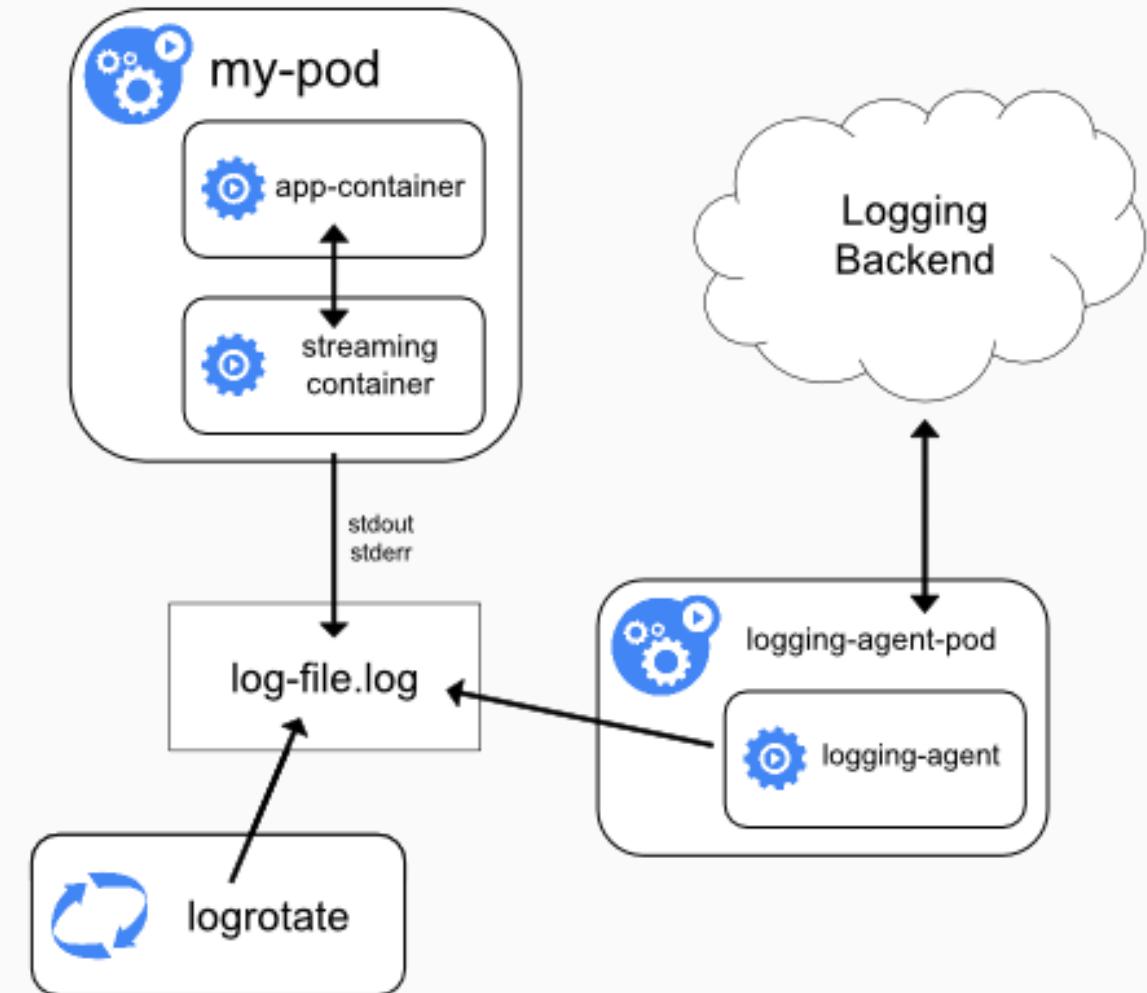


# Kubernetes - Logging



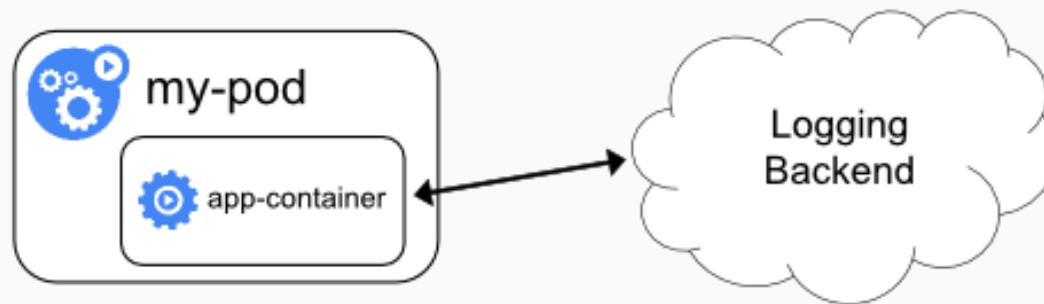
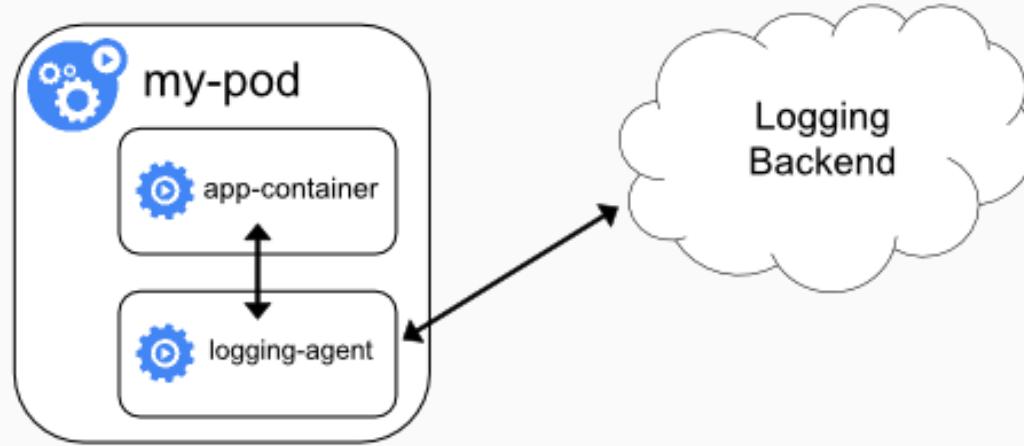
# Kubernetes

## - Logging



# Kubernetes

## - Logging



# Elasticsearch



elasticsearch

>> logging server

- Logs
- Metrics
- APM
- Uptime
- Site Search
- App Search
- Workspace Search
- Maps
- SIEM
- Endpoint Security

<https://www.elastic.co/elasticsearch/>

# Elasticsearch

So what is wrong?

# Elasticsearch

---

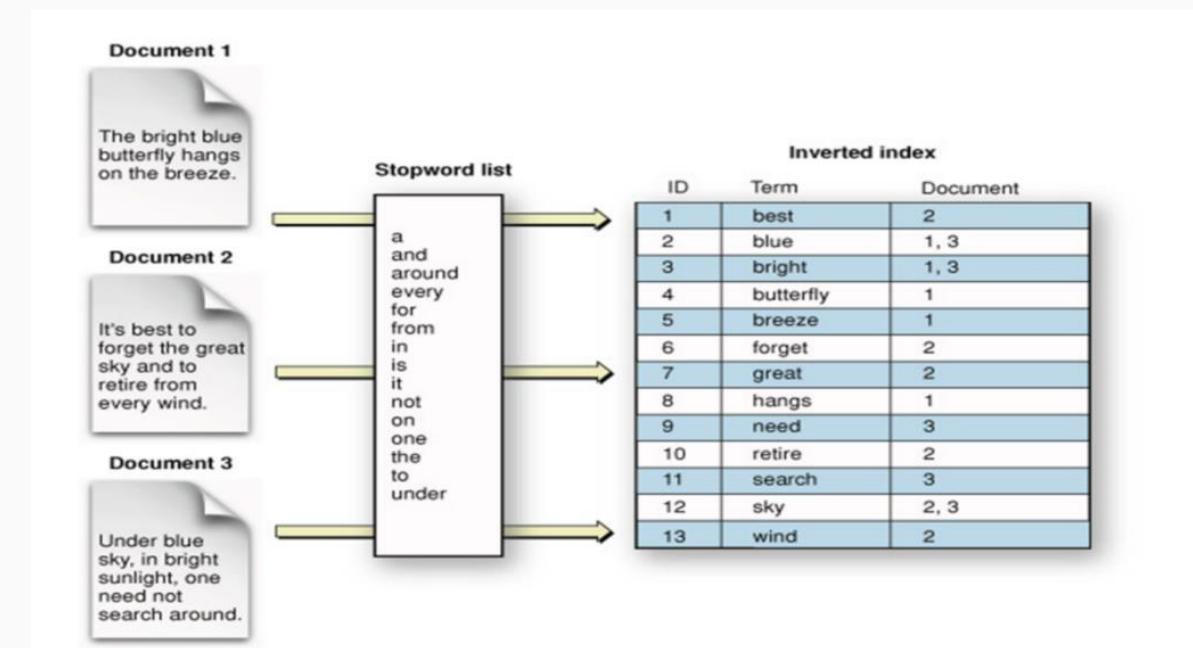
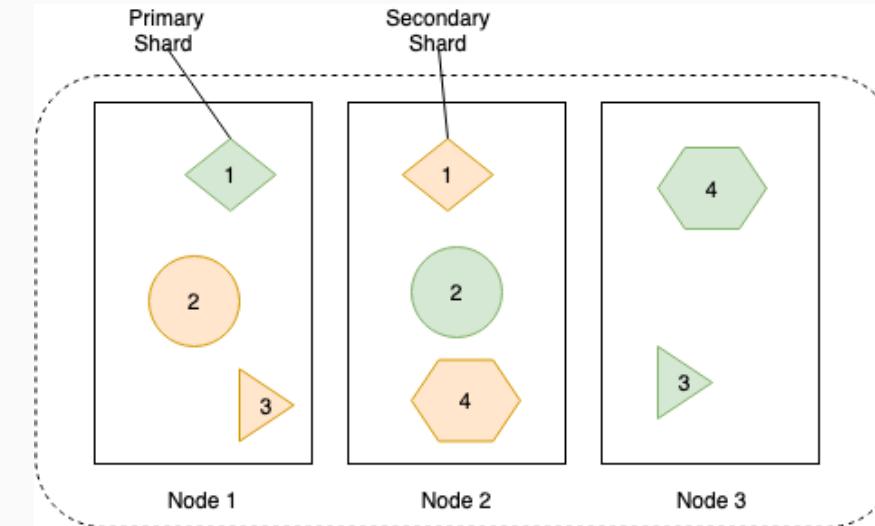
## Architecture

There are following type of nodes in the cluster:

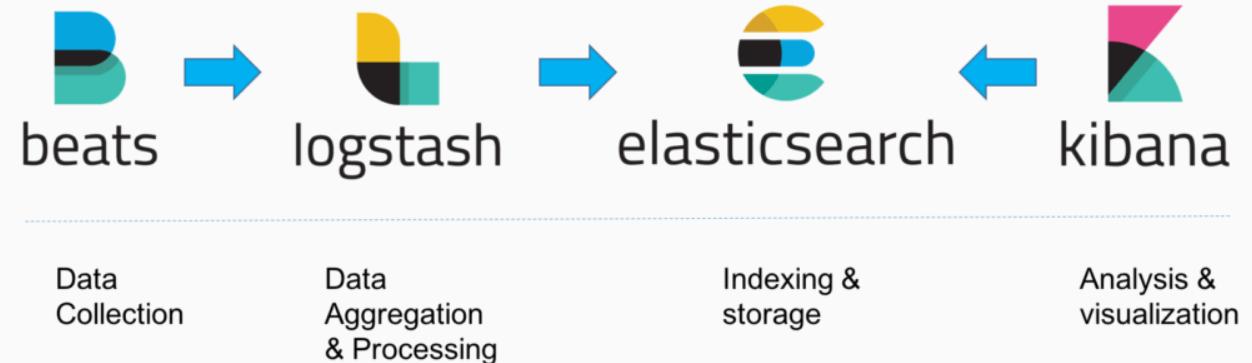
- Master Nodes – controls the cluster, requires a minimum of 3, one is active at all times
- Data Nodes – to hold index data and perform data-related tasks
- Ingest Nodes – used for ingest pipelines to transform and enrich the data before indexing
- Coordinating Nodes – to route requests, handle search reduce phase, coordinates bulk indexing
- Alerting Nodes – to run alerting jobs
- Machine Learning Nodes – to run machine learning jobs

# Elasticsearch

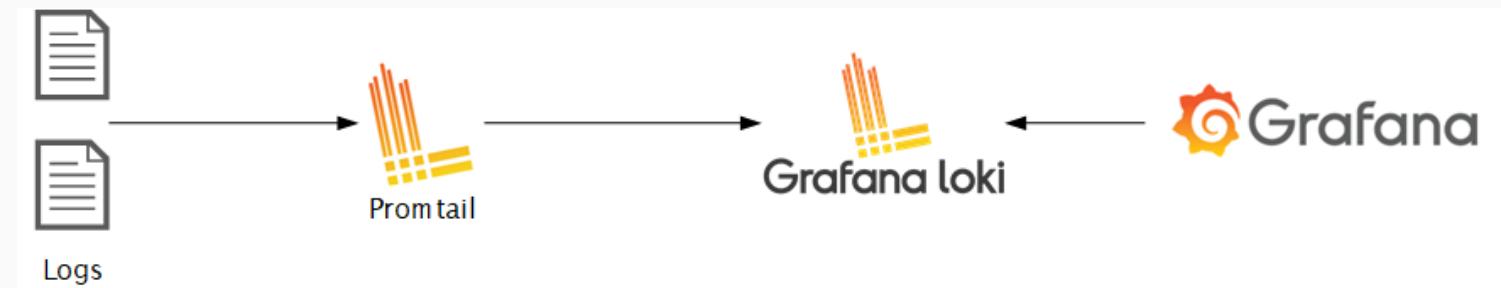
## Architecture



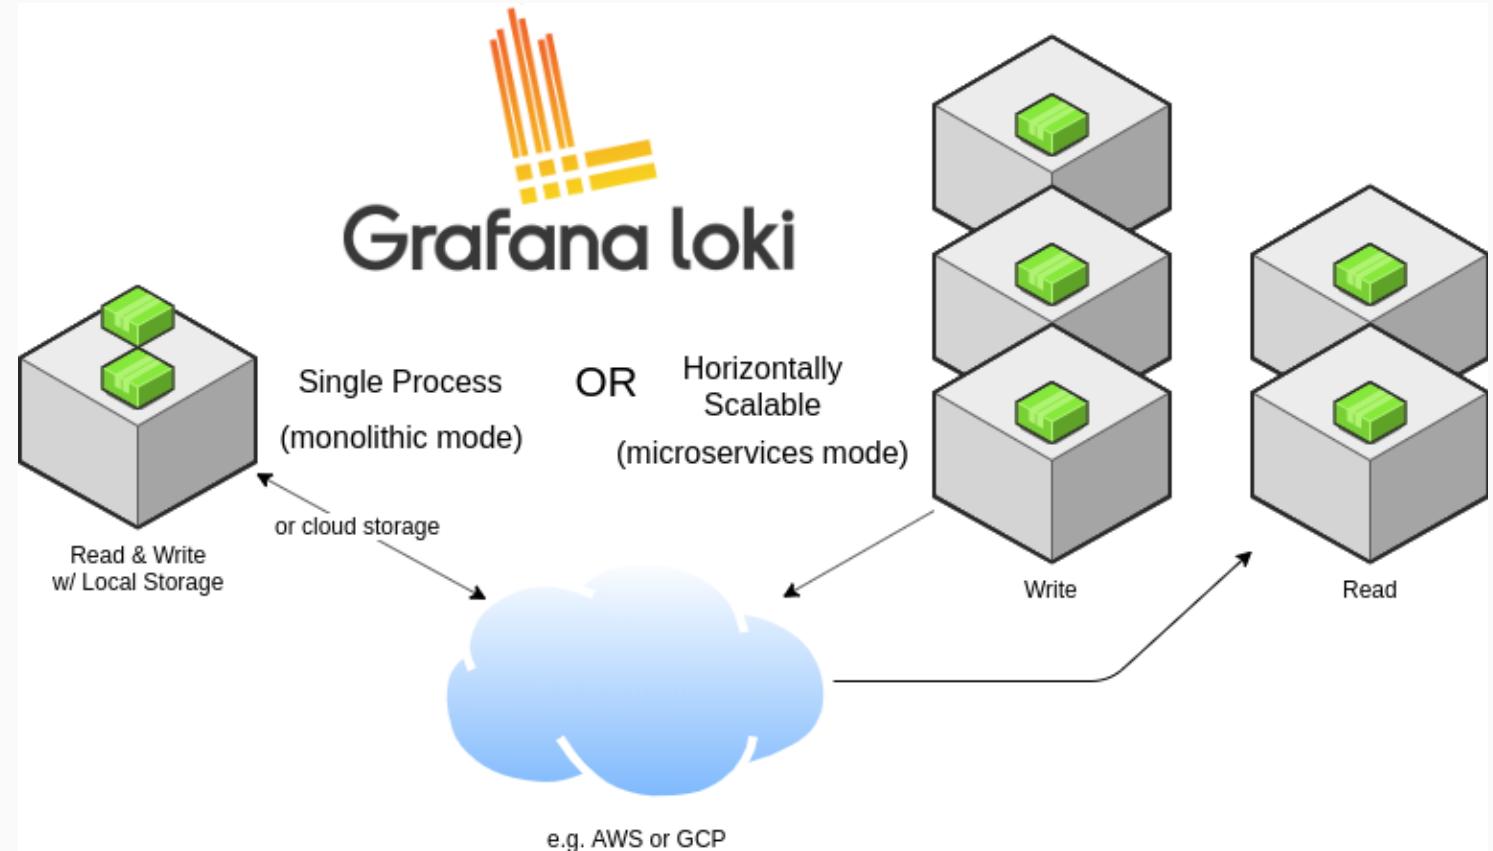
# ELK, EFK?



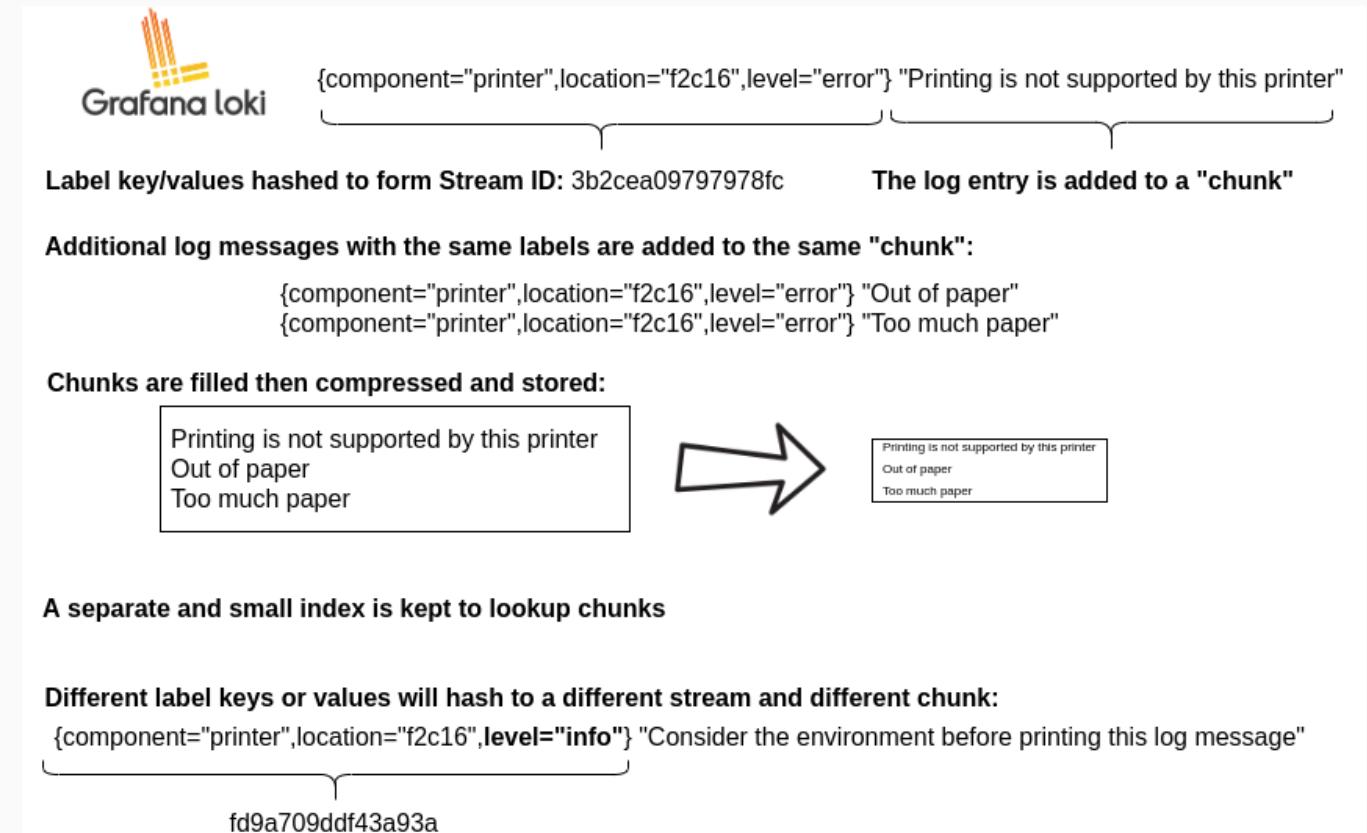
# LOKI



# LOKI - Architecture

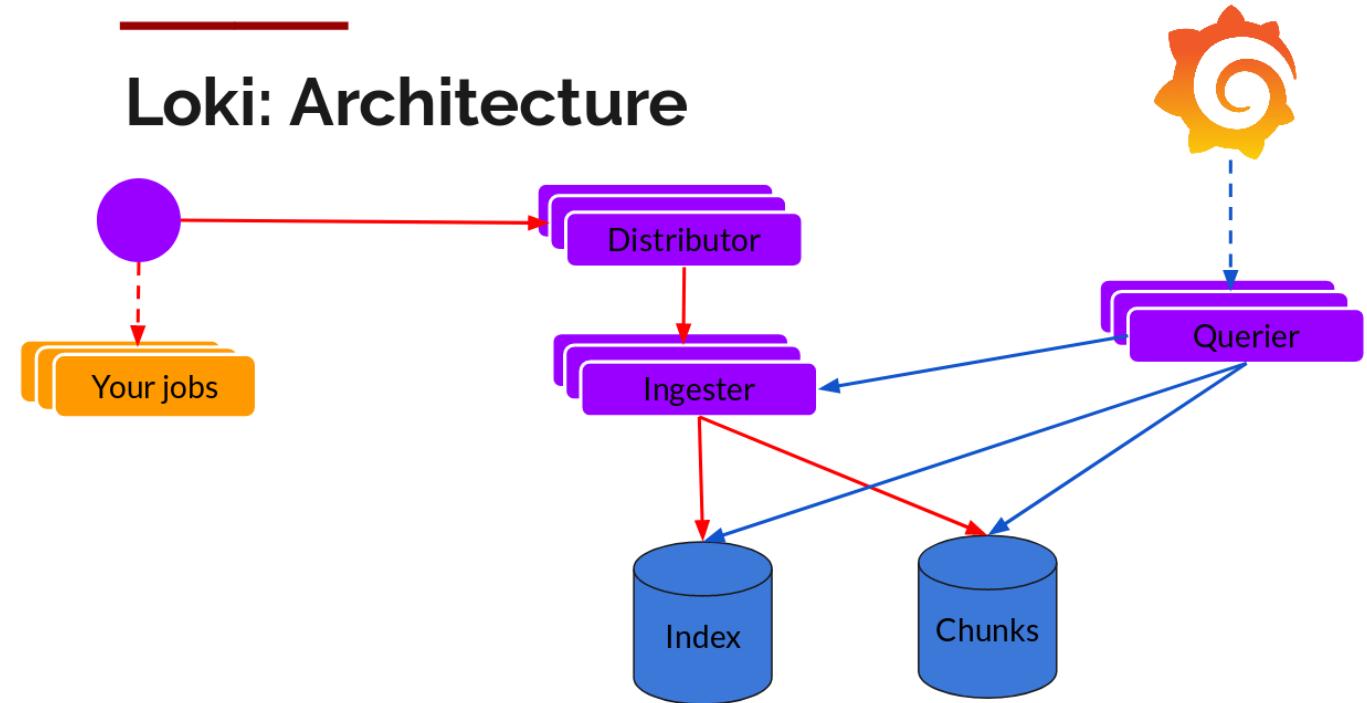


# LOKI - Architecture



# LOKI - Architecture

## Loki: Architecture



# **LOKI -**

# **LogQL**

<https://grafana.com/docs/loki/latest/logql/>

# Elasticsearch vs LOKI (only from logging perspective)

ITEM	Elasticsearch	Loki
<b>Query Language</b>	DSL and Lucene – very powerfull with extensive operator support	LogQL – inspired by Prometheus, basic operator support
<b>Scalability</b>	Support	Support – but better due to microservice architecture. You can easely seperate read from write. Can consume much bigger ammount of data.
<b>Multi-tenancy</b>	one index per tenant, tenant-based routing, using unique tenant fields, and use of search filters	X-Scope-OrgId in the HTTP header request
<b>Cost</b>	\$\$\$\$\$\$\$\$\$\$....	Extreamly cost effective

# Monitoring

# Monitoring

1. Po co monitorować?
2. Co monitorować?

**Monitoring –  
co  
monitorować?**

**Wszystko**

# Monitoring – co monitorować?

- Fizyczne elementy:
  - Serwery
  - Urządzenia sieciowe
- Sieć:
  - Dostępność (z zewnątrz)
  - Opóźnienia/straty/jitter
- OS:
  - HDD
  - CPU
  - RAM, itd.
- Aplikacje:
  - Czy jest uruchomiona
  - Czy odpowiada
  - Czy działa
  - Metryki
  - Logi
  - Performance
- Trendy

# Monitoring – co monitorować?

- Fizyczne elementy:
  - Serwery
  - Urządzenia sieciowe
- Sieć:
  - Dostępność (z zewnątrz)
  - Opóźnienia/straty/jitter
- OS:
  - HDD
  - CPU
  - RAM, itd.
- Aplikacje:
  - Czy jest uruchomiona
  - Czy odpowiada
  - **Czy działa**
  - **Metryki**
  - **Logi**
  - **Performance**
- **Trendy**

# **Monitoring – co monitorować?**

A co jeżeli nie wiemy co monitorować w aplikacji?  
**RED**

# Monitoring - sposoby

**Ze względu na sposób:**

- Pull
- Push

**Ze względu na miejsce:**

- Centralnie
- Rozproszenie (np. poprzez agentów)

# Monitoring - narzędzia

## **Open-source (legacy):**

- Zabbix
- Nagios
- Icinga
- Sensu

## **Open-source (new):**

- Prometheus
- Jaeger/Tempo
- Telepresence

## **SaaS:**

- DataDog
- NewRelic

Monitoring -  
Zabbix



# Monitoring - Zabbix

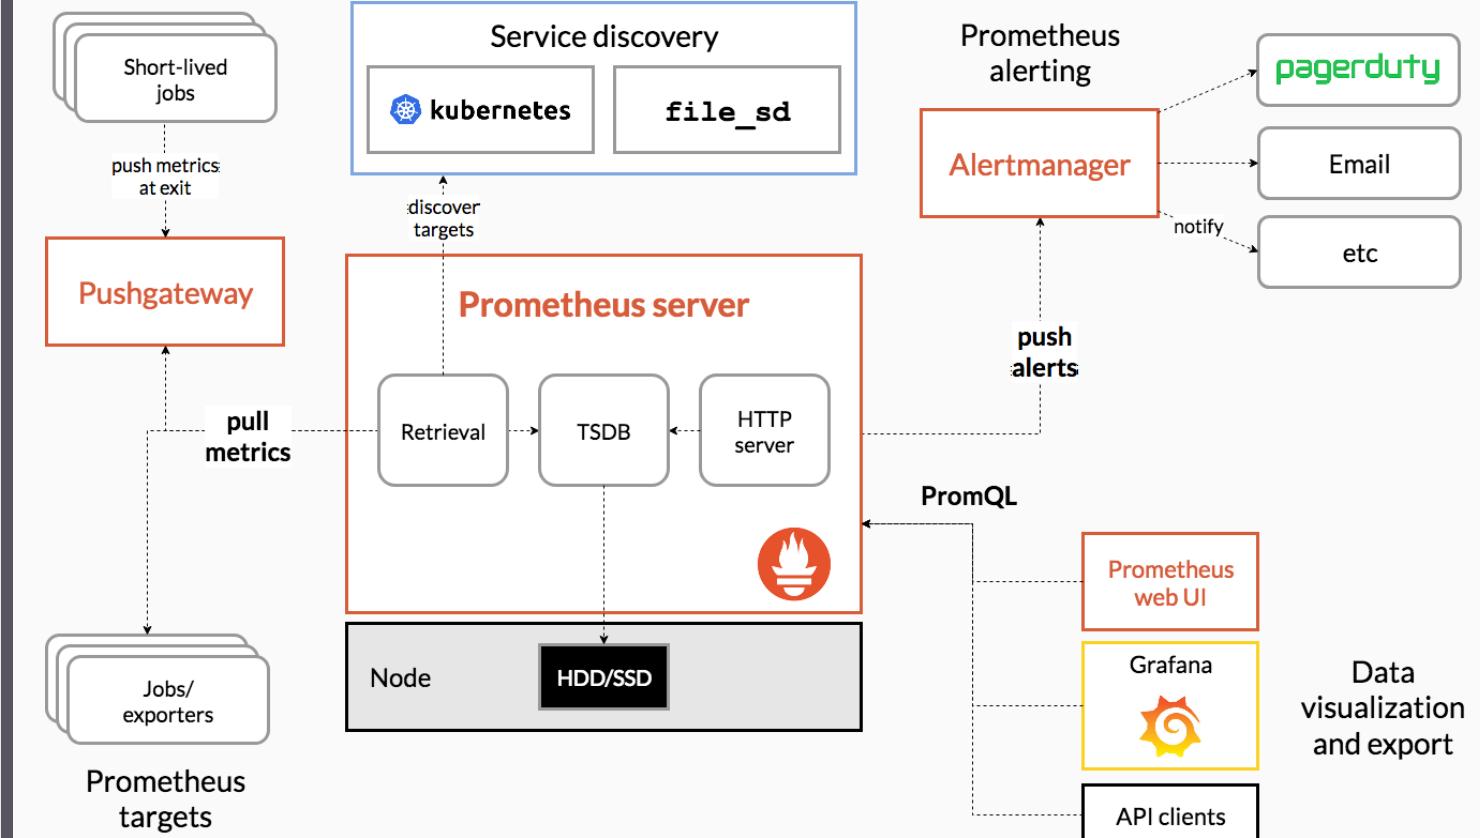
## Why Zabbix?

- Data gathering
- Highly configurable alerting
- Web monitoring capabilities
- Extensive visualization options
- Historical data storage
- Easy configuration
- Use of templates
- Network discovery
- Fast web interface
- Zabbix API
- Permissions system
- Full featured and easily extensible agent
- Binary daemons
- Ready for complex environments

# **Monitoring – Zabbix – Installation**

<https://www.zabbix.com/documentation/current/manual/installation>

# Monitoring – Prometheus – Architecture



# Monitoring – Prometheus – endpoints

```
<metric name>{<label name>=<label value>, ...}
```

```
api_http_requests_total{method="POST", handler="/messages"}
```

# Monitoring – Prometheus – querying

<https://prometheus.io/docs/prometheus/latest/querying/basics/>

# Monitoring – Prometheus – important topics to consider

- Security
- Retention period (default 15d)

<https://prometheus.io/docs/prometheus/latest/storage/>

# The 12 Factor App

## Design Considerations in the World of Containers

### EMERGING STANDARD

Platform and language agnostic app will be app architecture. Deployed, developed by teams and portably well suited to cloud platforms. This is one way to develop cloud native applications, but not the only way.

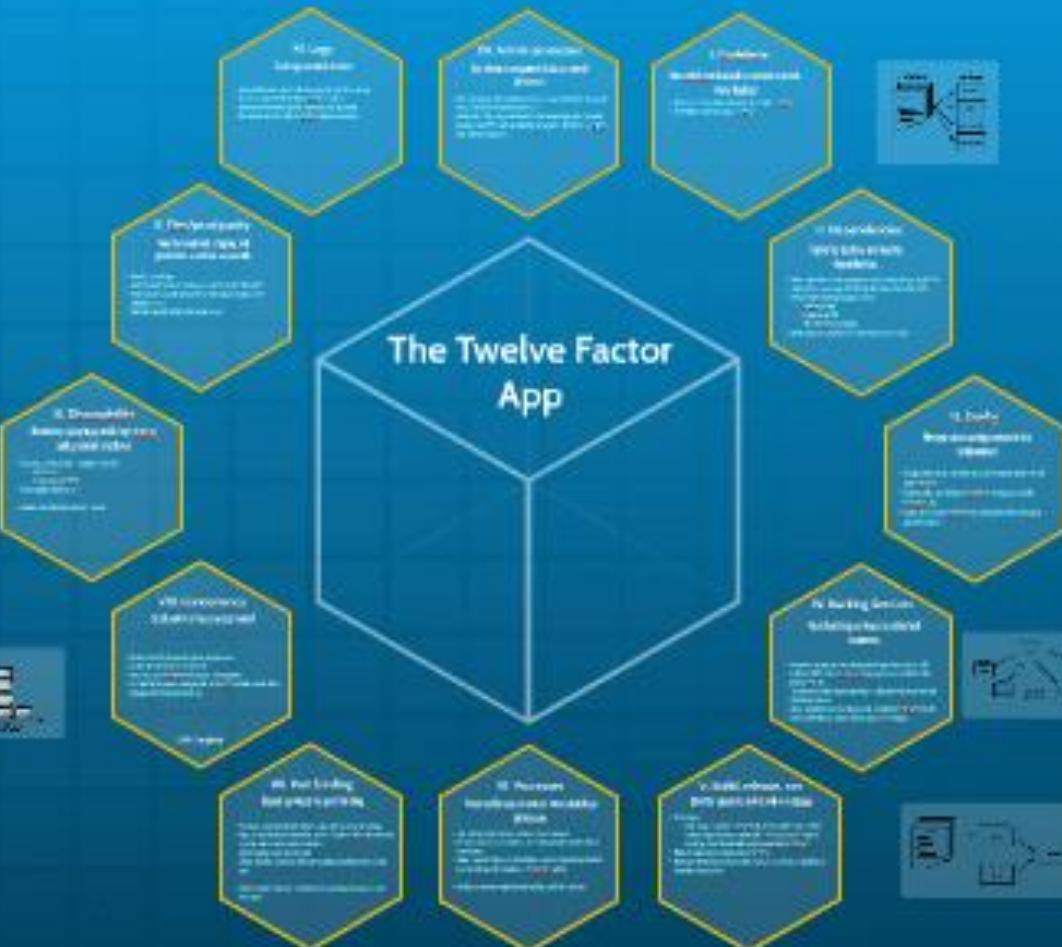
#### Example:

- declarative formats
- user contract with CI for more portability
- describes its cloud platform
- removes divergence between Dev & Prod
- automated scaling

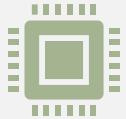
### Cloud Native vs. Container



## The Twelve Factor App



# 12th factor app for SaaS applications



Agnostic? What is that?



What is Software as a Service



Declarative approach



Contract

**1 application**  
**1 codebase**  
**many**  
**Deploys**



Many codebases means many applications



Common codebase means lack of dependency manager



Deploy from one codebase to each env (dev/stage/prod)

# Explicitly declare and isolate dependencies



Each application declares all dependencies



Application is not dependent on „not listed” dependencies



Application is not dependent on system tools (like bash/ps/curl)

# Store config in the environment



Resource handles to the database



Credentials to external services such as Amazon S3 or Twitter



Per-deploy values such as the canonical hostname for the deploy



Can I share my app as open source right away without compromising information?



Can you run my app on your environment using just code?



Use ENV variables to set configuration

# Treat backing services as attached resources

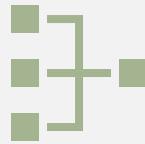


All services are attached, there is no 'local' and 'remote' resources



Service can be easily attached and detached

# Strictly separate build and run stages



Build code (merge all dependencies)



Prepare release with version  
(use build and config)



Run release – should be possible without supervision

# Execute the app as one or more stateless processes



Any data that needs to persist should be attached through resource



Application can't be sure that information on disk or in memory would be available



Move „sticky sessions” to REDIS or similar tool

# Export services via port binding



Attach PORT to each possible application



If application doesn't use PORT then it is not service (job/run proces)

# Scale out via the process model



Consider that application will be scaled horizontally by increasing number of processes



Don't start process in background

# Maximize robustness with fast startup and graceful shutdown



processes are disposable – they can be stopped at any time



Handle SIGTERM gracefully



Consider „power off“ situations

# Keep development, staging, and production as similar as possible



developers who wrote code are closely involved in deploying it and watching its behavior in production.



**The twelve-factor developer resists the urge to use different backing services between development and production**

# Treat logs as event streams



A twelve-factor app never concerns itself with routing or storage of its output stream. It should not attempt to write to or manage logfiles.

# Run admin/ma gement tasks as one-off processes

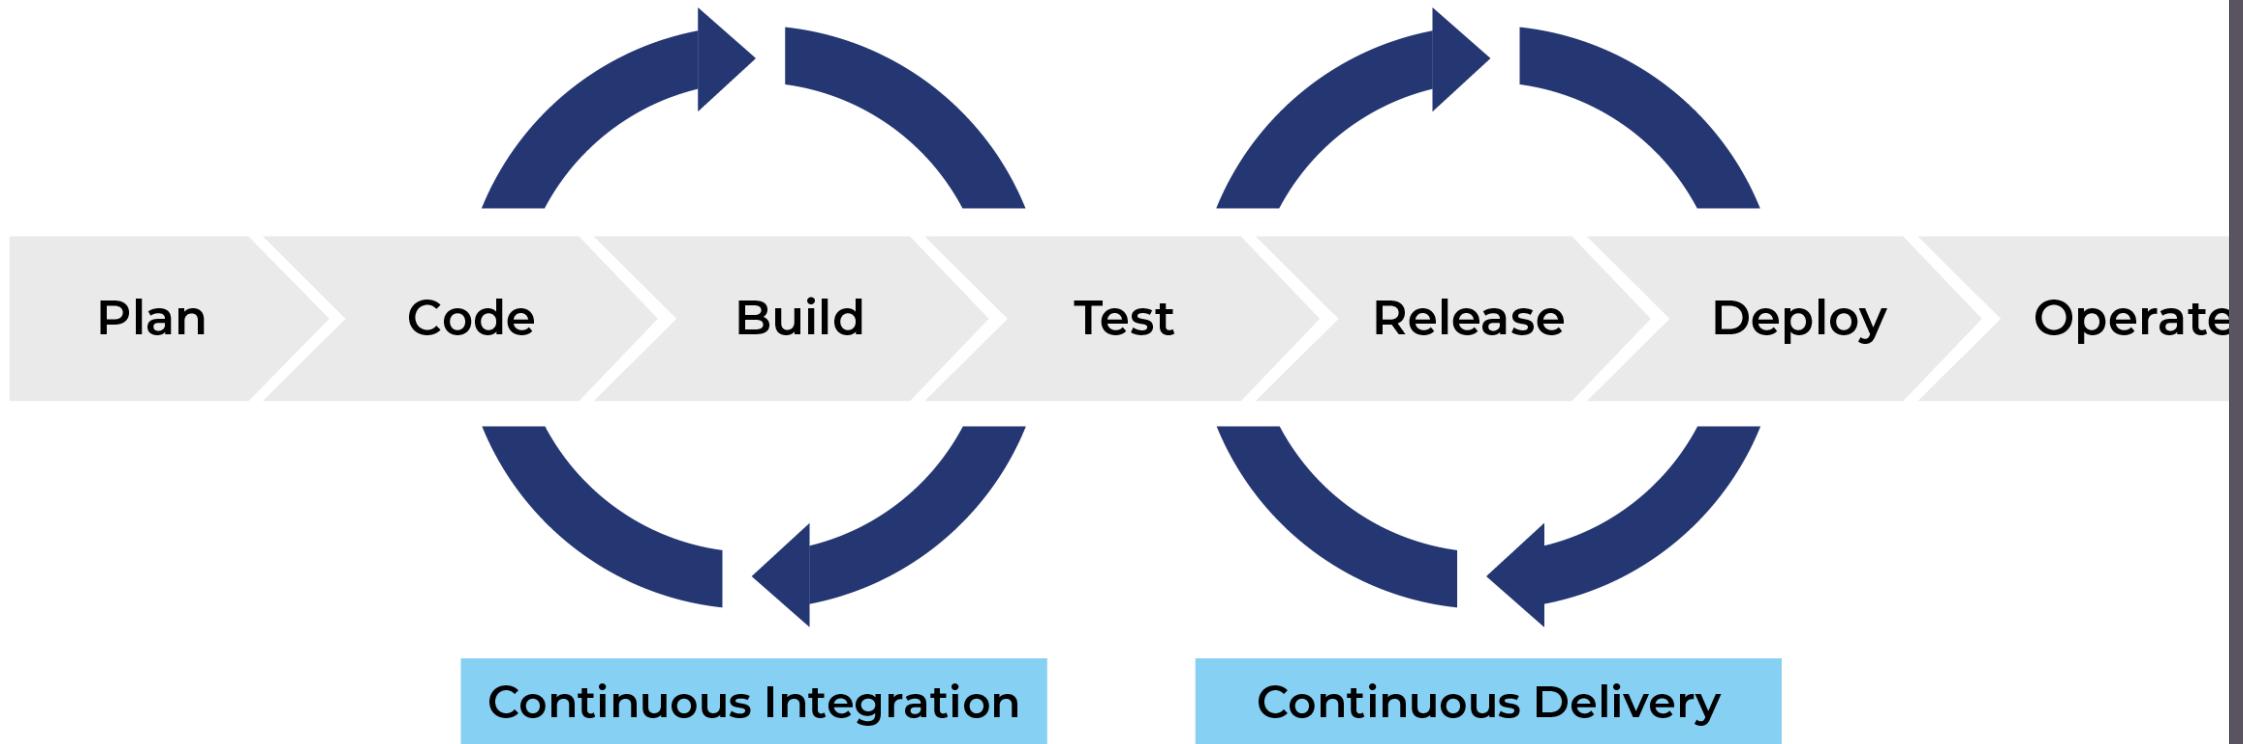


One-off admin processes should be run in an identical environment as the regular long-running processes of the app. They run against a release, using the same codebase and config as any process run against that release.

# Documentation

<https://12factor.net>

# CI/CD



# CI/CD

## Pipeline



Split State from  
Code

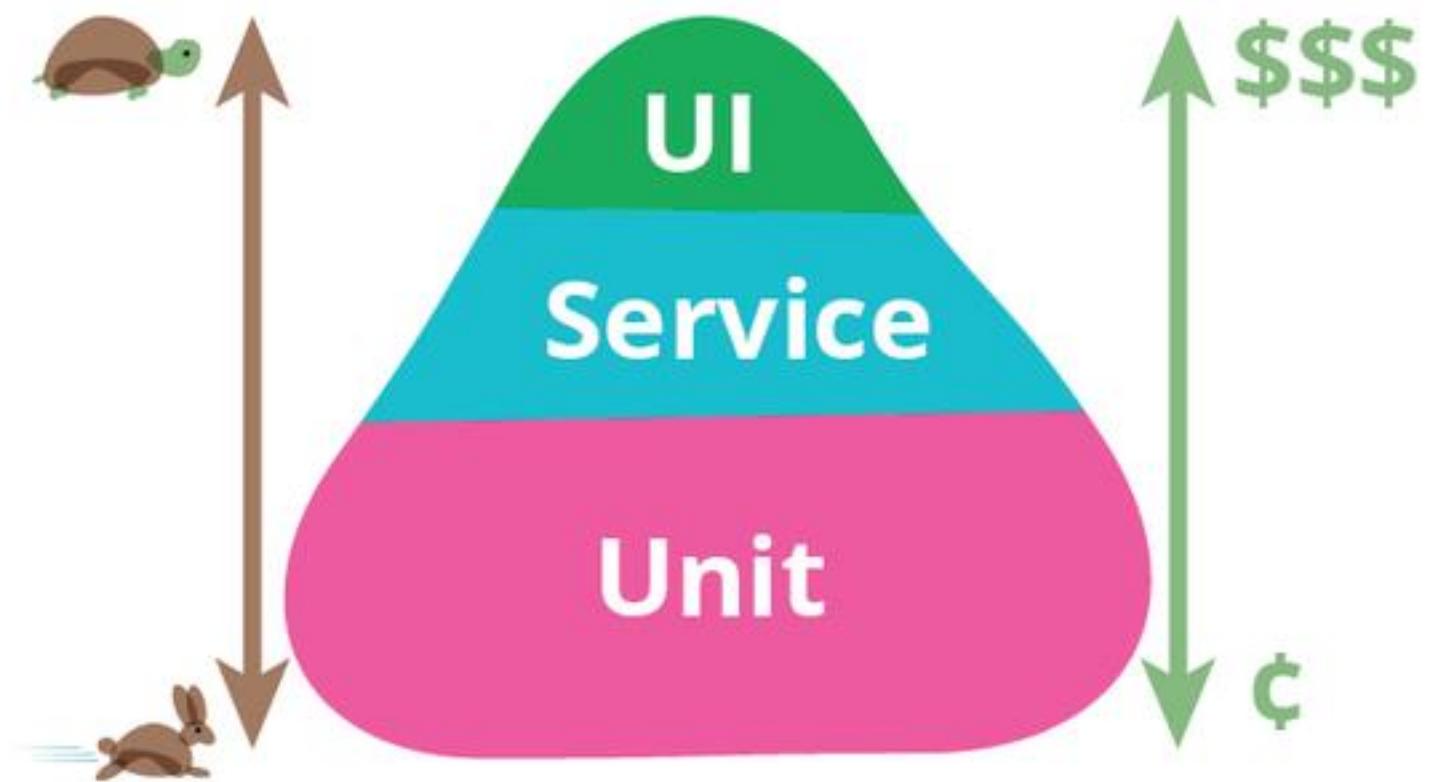
Code repository  
State repository  
(k8s yaml)



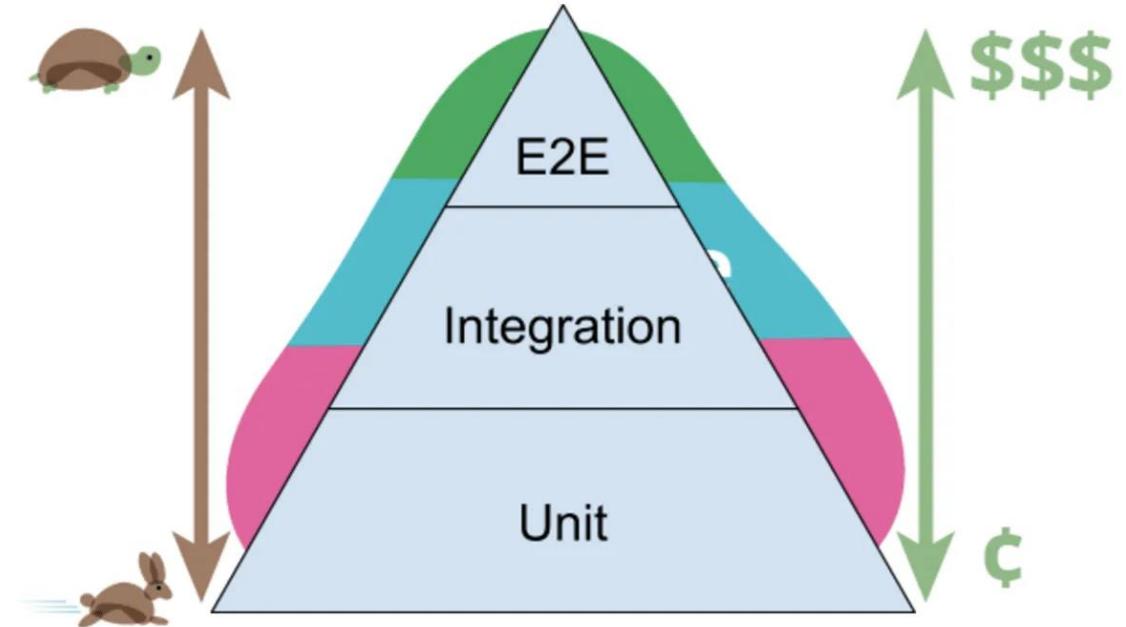
Master is the  
TRUTH

Disable push to  
master (only PR)  
Consider  
yamllint  
Place pipeline  
into repository

# CI/CD Pipline



# CI/CD Pipline



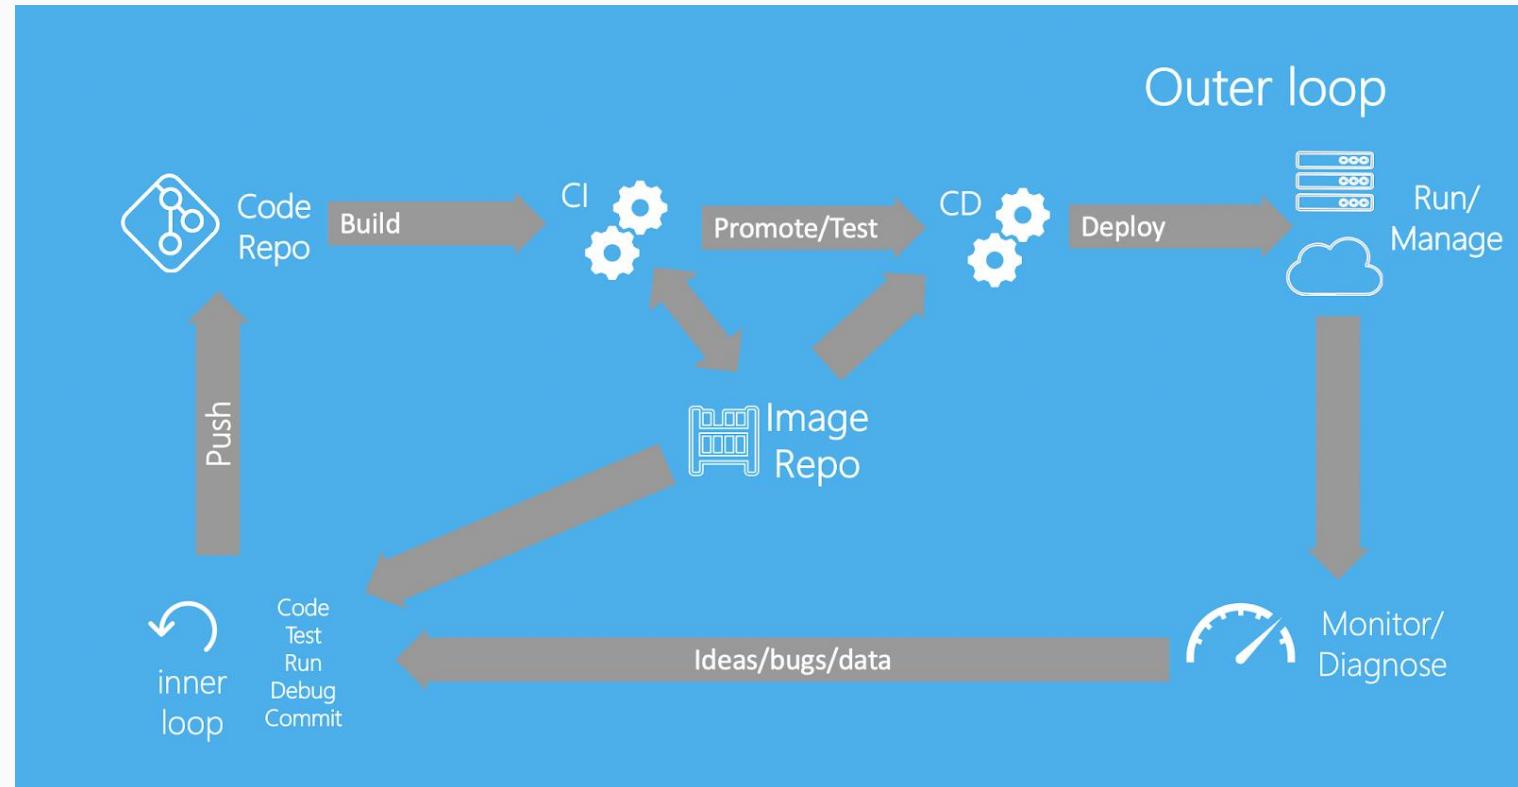
# CI/CD Pipline



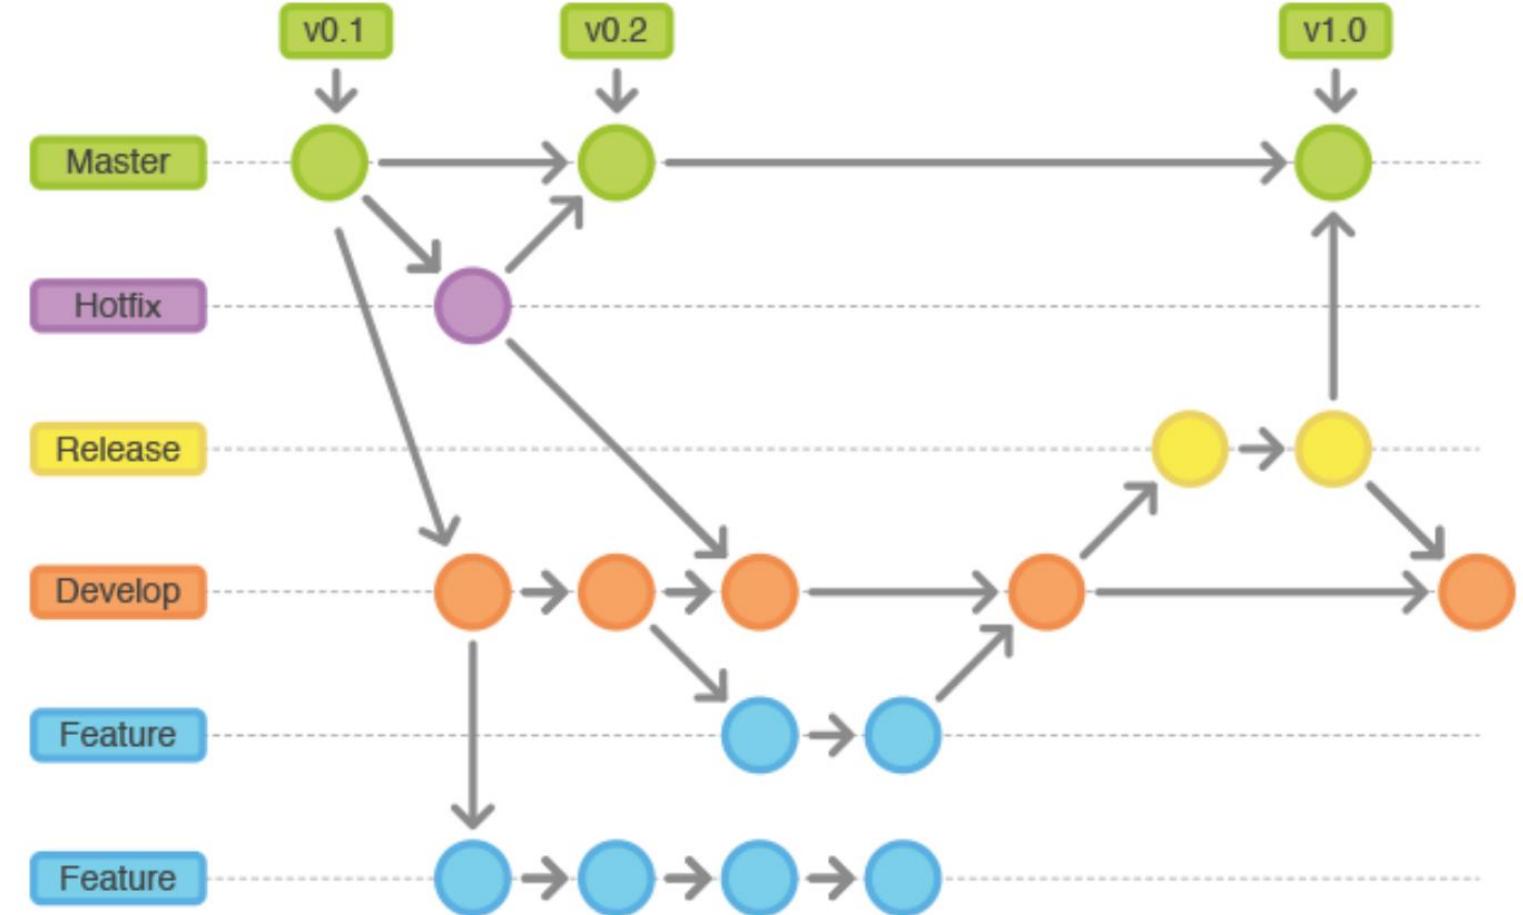
**CI/CD**  
**Pipeline**

**CI vs CD?**

# CI/CD Pipeline



# CI/CD Pipeline



# CI/CD

# Pipeline



HELM / Kustomize



Consider only templating



Kubectl apply -prune -namespace mobees -l app=mob



Own images



Own templates



Security

**CI/CD**

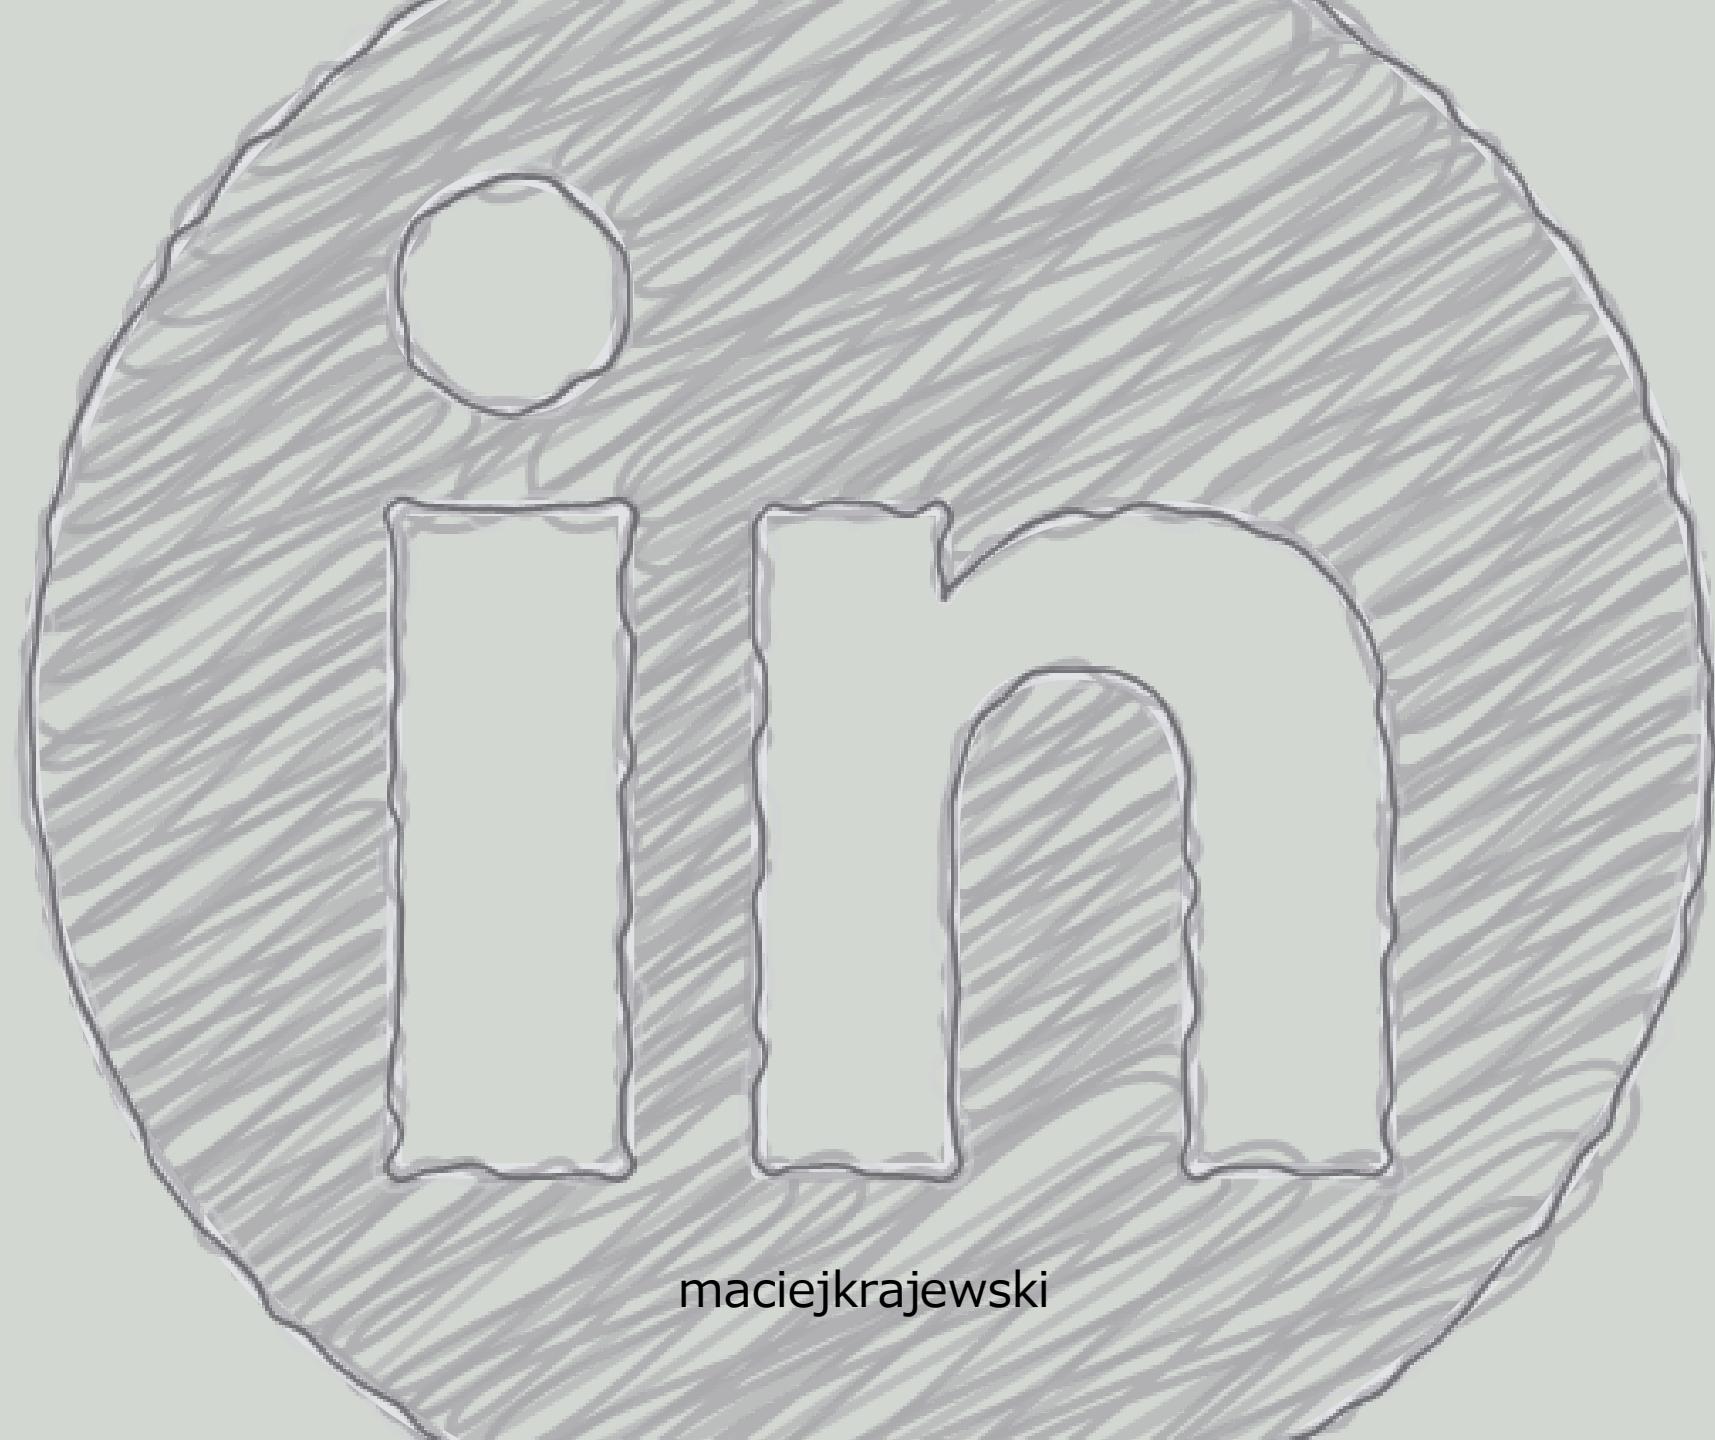
**HELM**

**&**

**KUSTOMIZE**



***kustomize.io***



**Dziękuję  
za uwagę**

maciejkrajewski